

# CS 221: Artificial Intelligence

## Planning (and Basic Logic)

Peter Norvig and Sebastian Thrun

Slide credits: Stuart Russell, Rina Dechter,  
Rao Kambhampati

# AI: Dealing with Complexity

---

- Agent Design
  - Reflex → Memory-Based
  - Reflex → Goal-Based → Utility-Based
- Environment
  - → Partially Observable, Stochastic, Dynamic, Multi-agent, Adversarial
- Representation
  - Atomic → Factored → Structured

# Finding Actions

---

- We've done: **Problem-Solving**

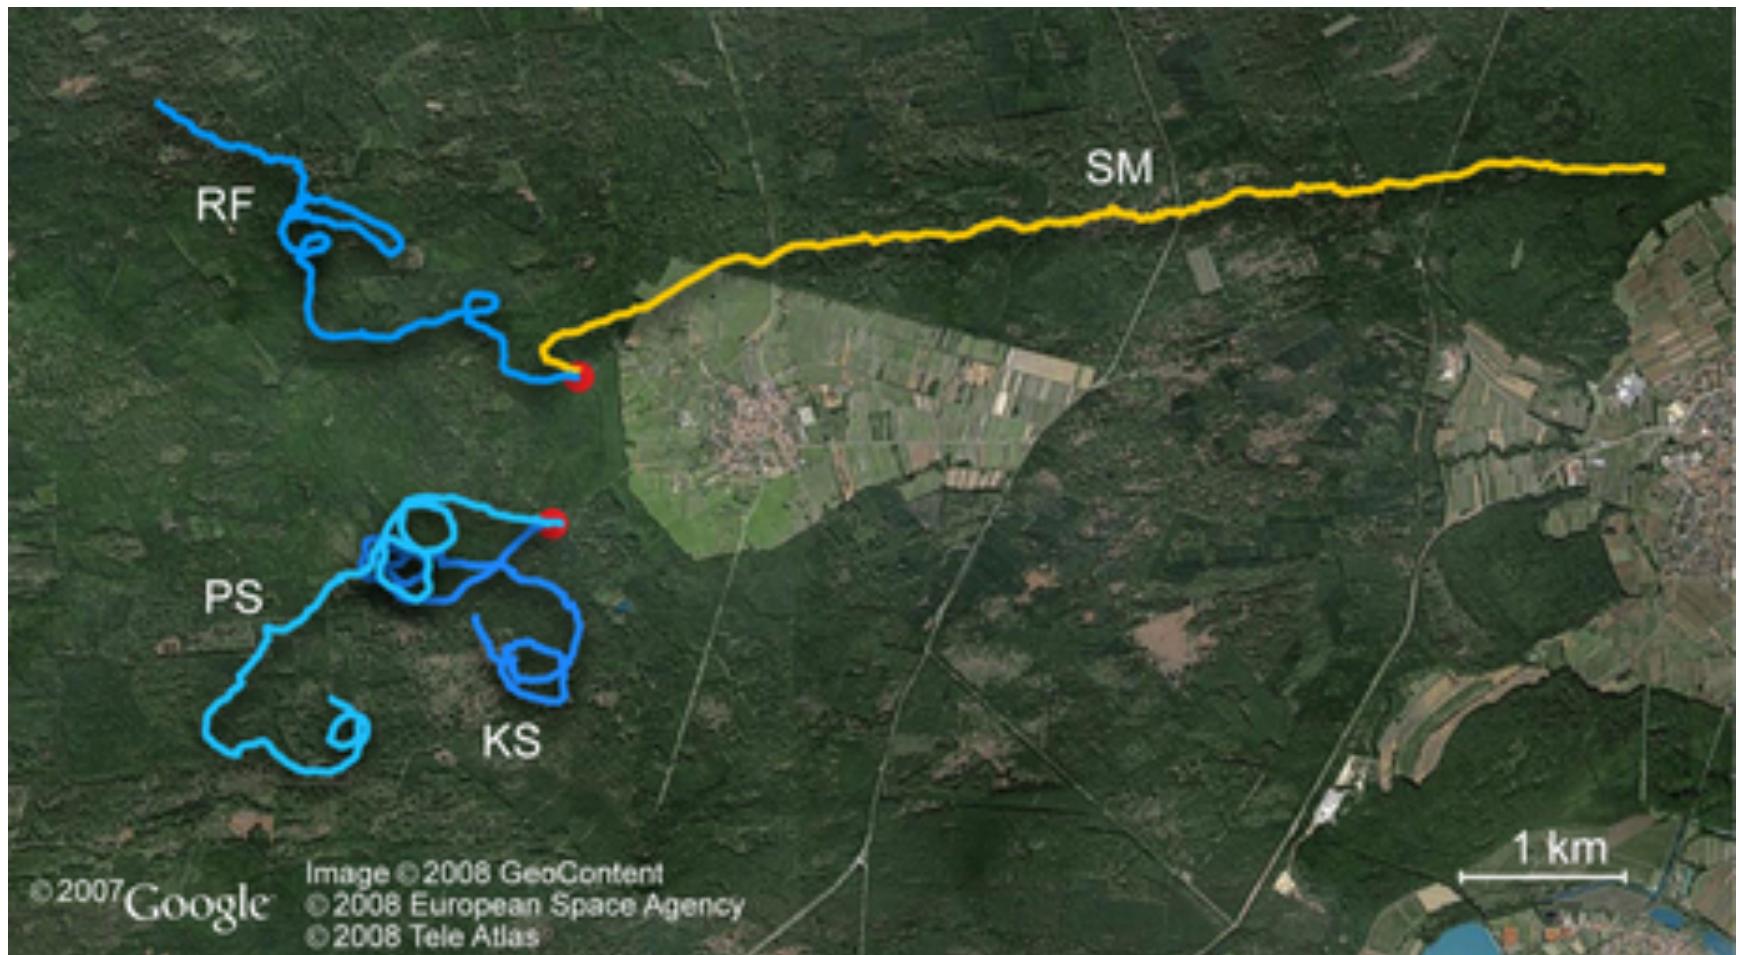
[Arad, Sibiu, Fagras, Bucharest]

- Introduce today: **Planning**

[while DoorLocked: try another key;  
push door; walk ahead]

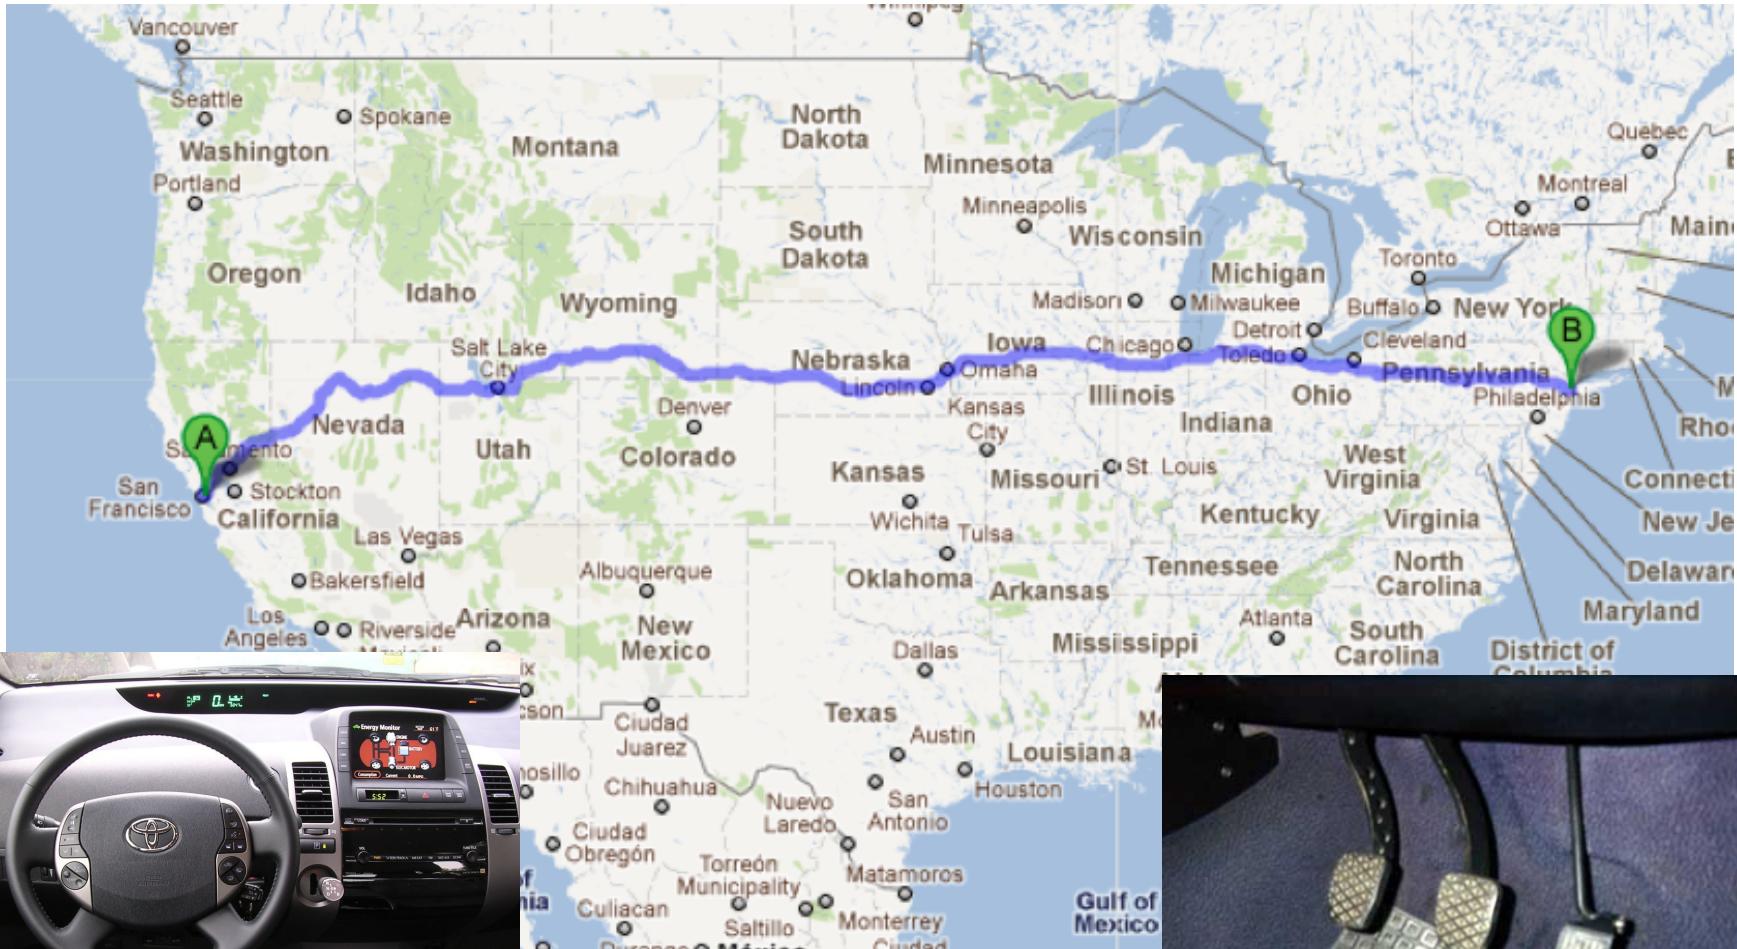
# What's wrong with Problem-Solving

---



Plan: [Forward, Forward, Forward, ...]

# What's wrong with Problem-Solving

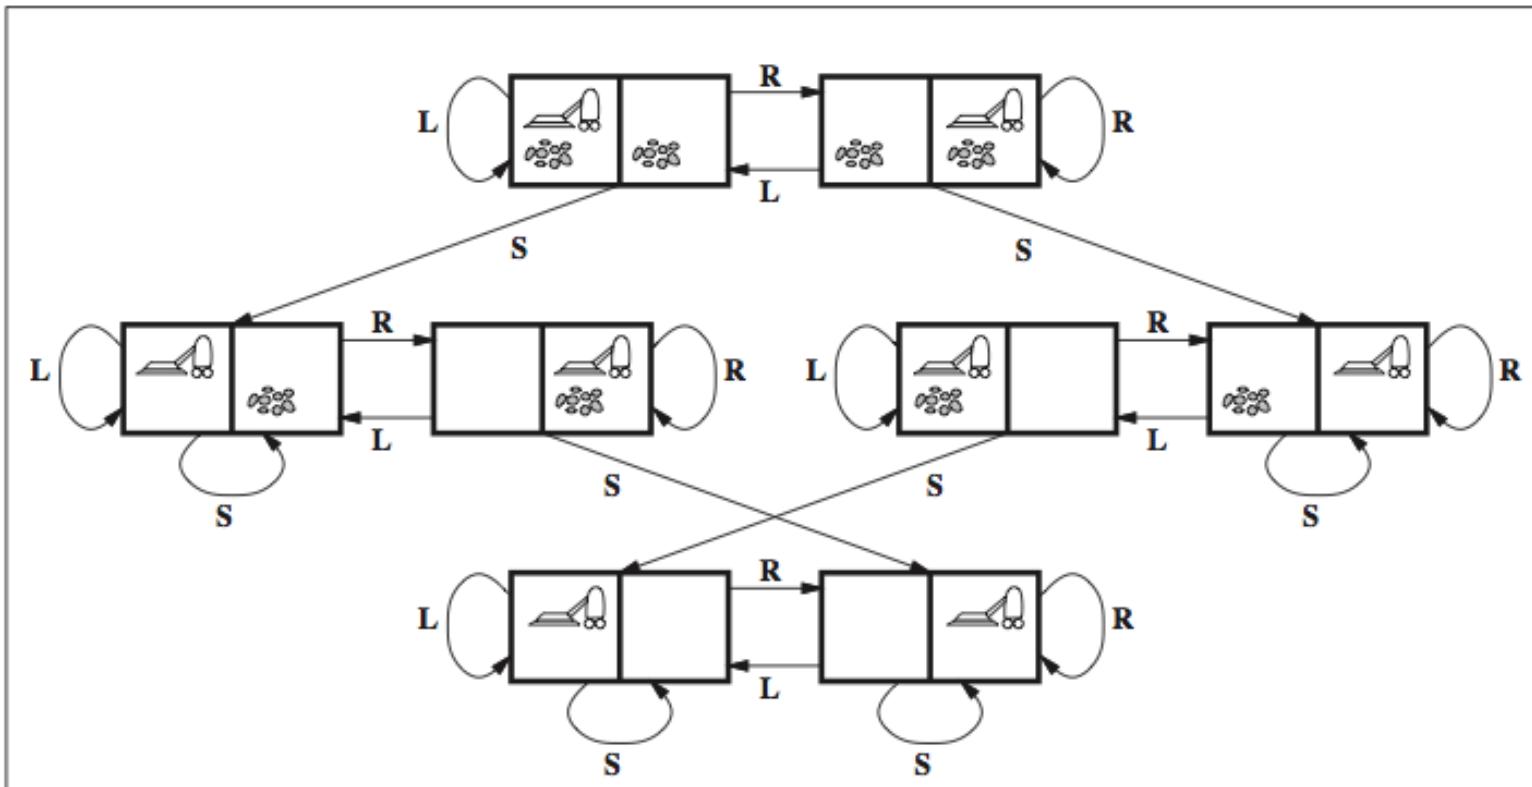


# Planning

---

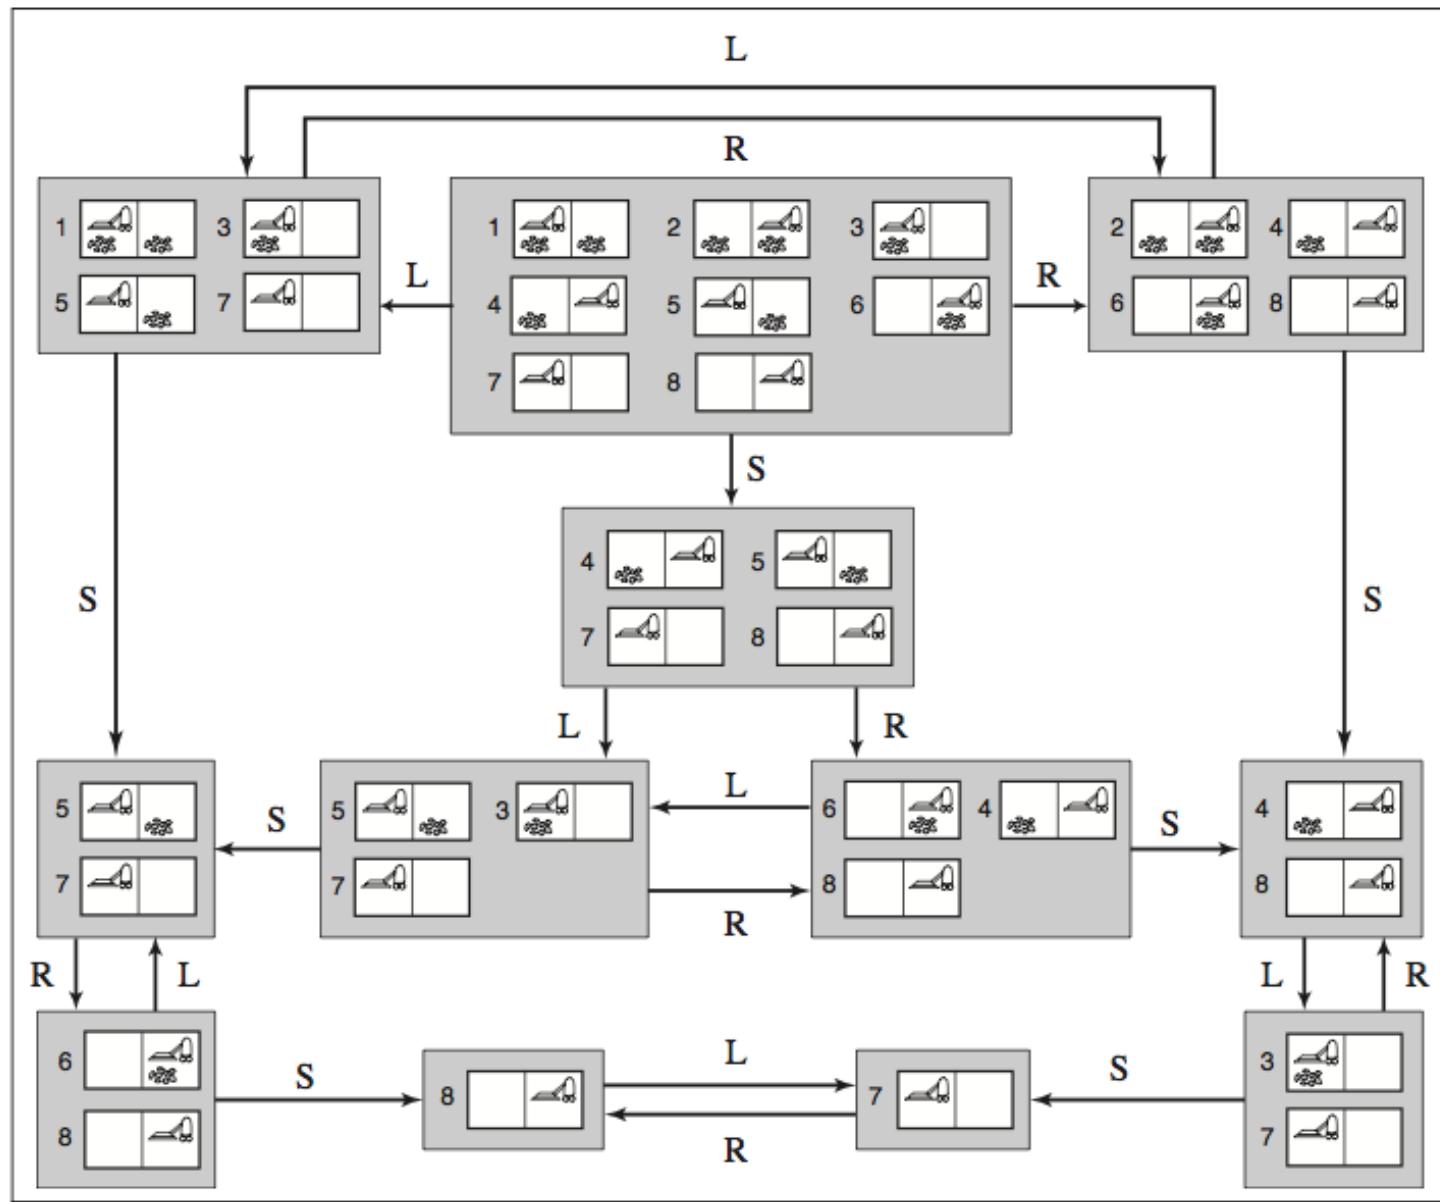
- A plan is a **program**
  - Not just a straight-line sequence of actions
- Planning and acting can be **interleaved**
  - Closed loop, not closed eyes
- Representation is more **flexible**
  - Can deal with partially-described states
  - Can deal with features, objects, relations
  - Can be hierarchical

# Dealing with Partial Observability: World vs. Belief States



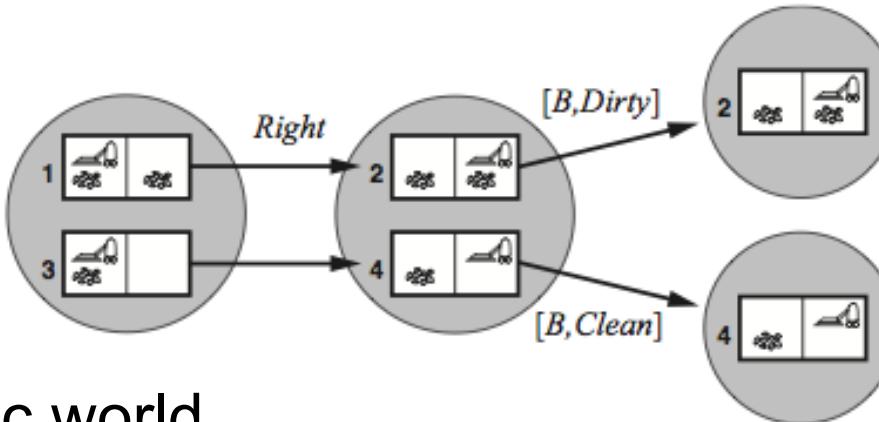
**Figure 3.3** The state space for the vacuum world. Links denote actions: L = *Left*, R = *Right*, S = *Suck*.

# Sensorless – Belief States - Conformant Plans



# Partial (local) Observability and Stochastic Worlds

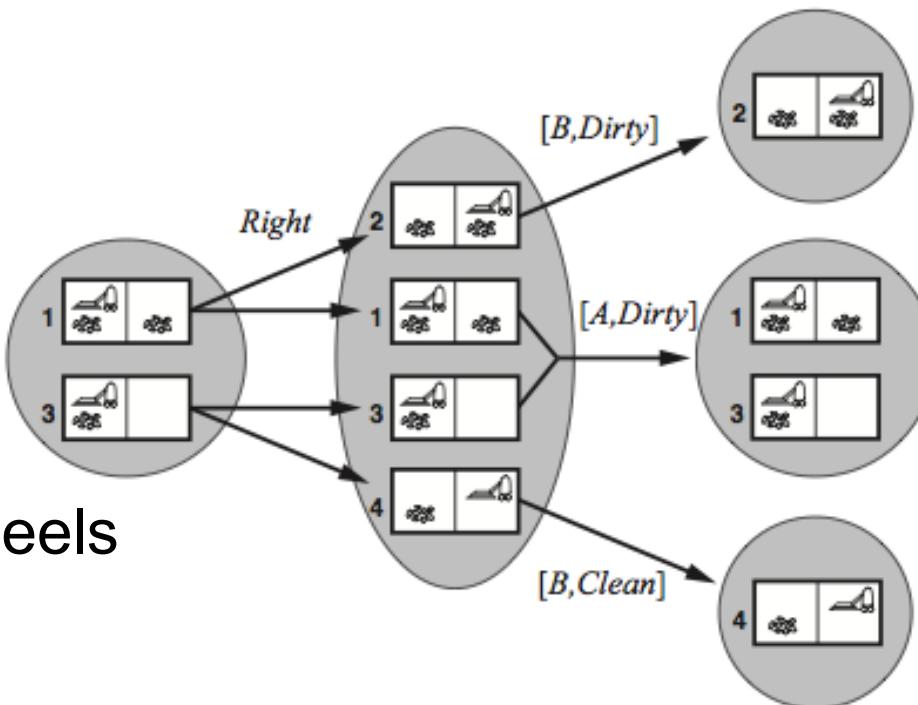
(a)



Deterministic actions; observe only local square

Deterministic world

(b)

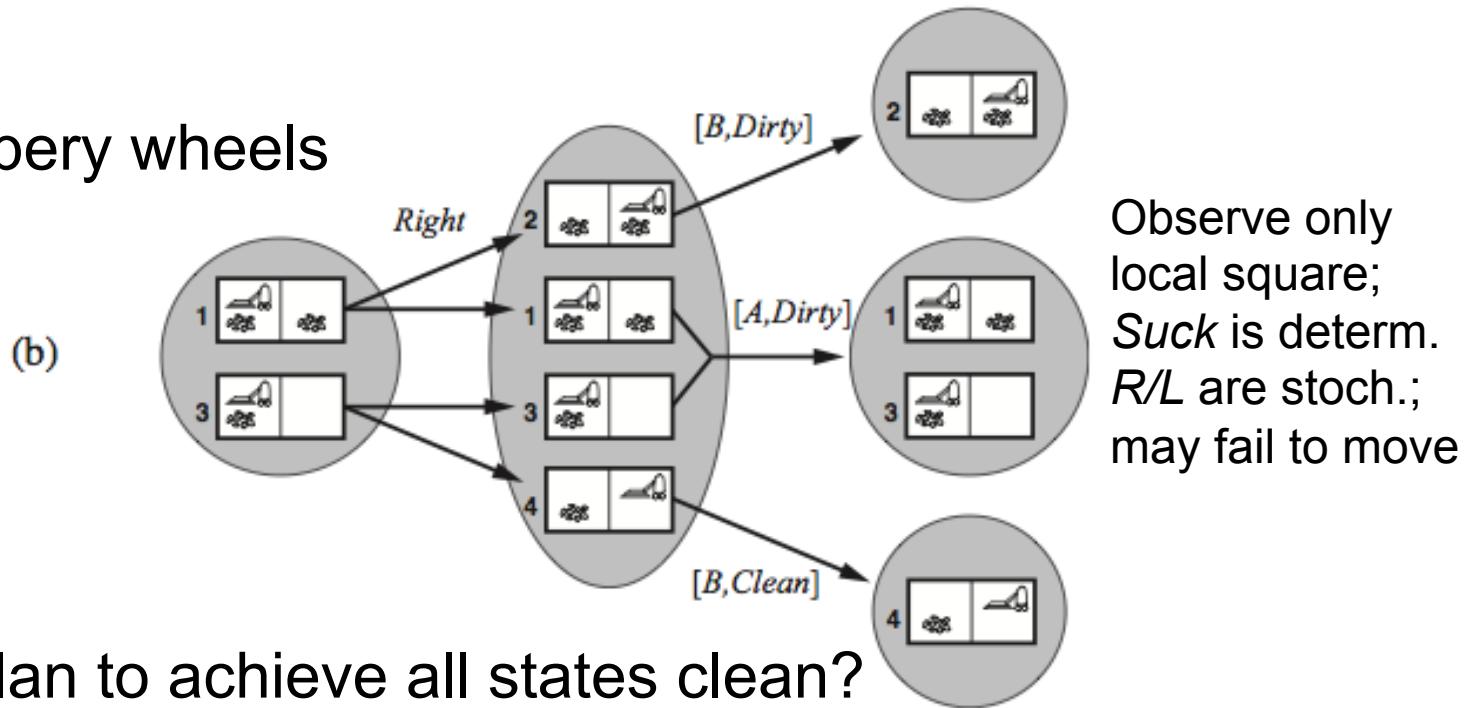


Observe only local square;  
Suck is determ.  
R/L are stoch.  
(may fail to move)

Slippery wheels

# Planning and Sensing in Partially Observable and Stochastic World

Slippery wheels



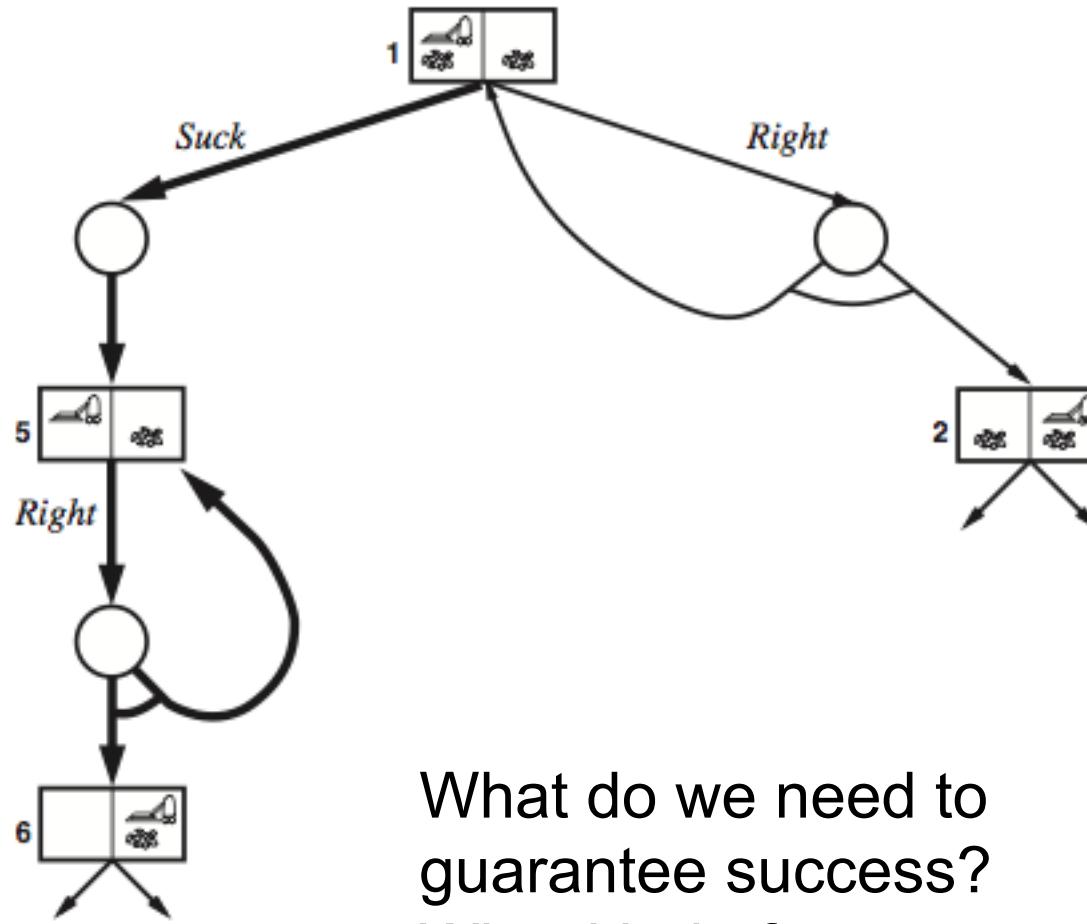
What is a plan to achieve all states clean?

[1:Suck; 2:Right; (**if A:** **goto** 2); 3:Suck]

also written as

[Suck; (**while A:** Right); Suck]

# Search Graph as And/Or Tree



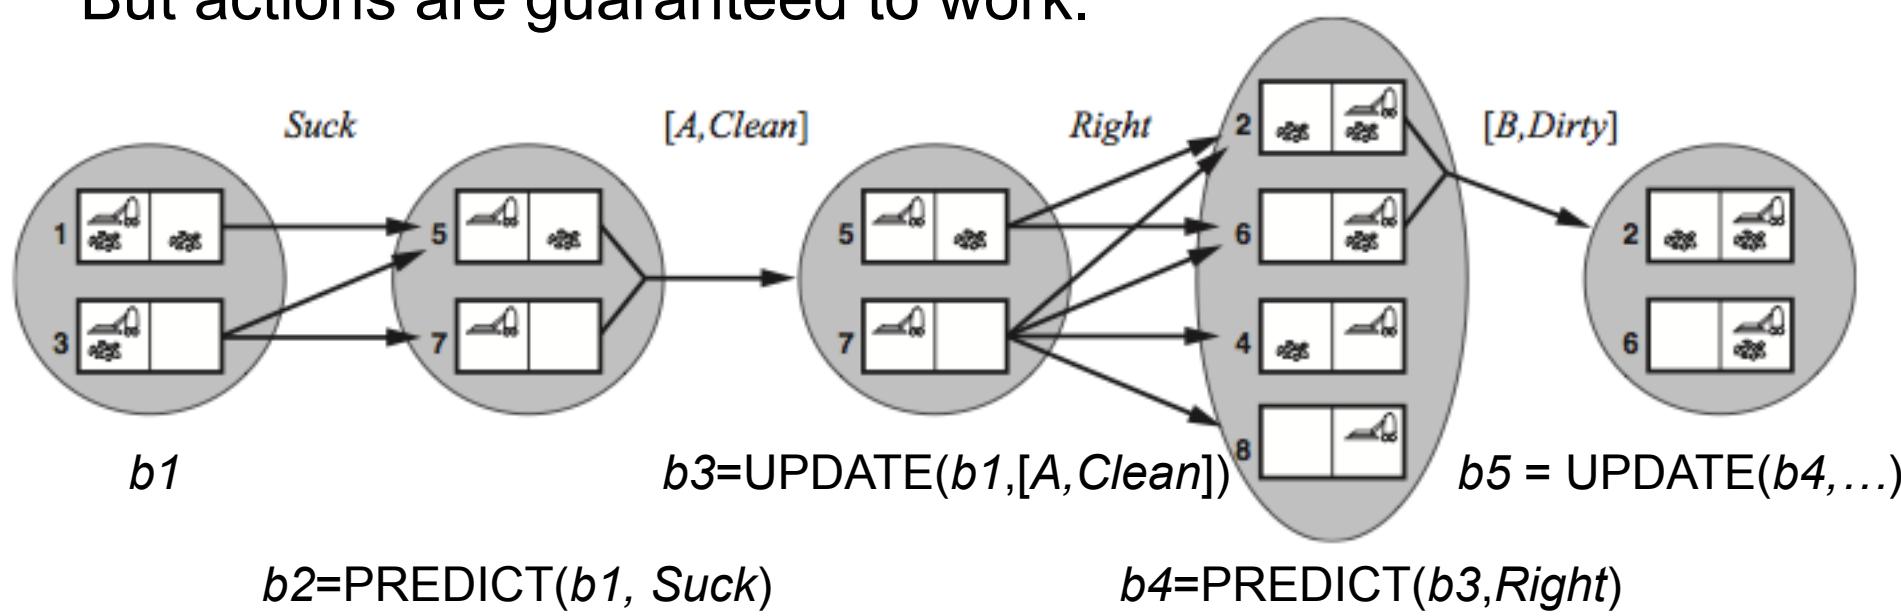
What do we need to  
guarantee success?  
What kind of guarantee?

# As Equations, not Tree

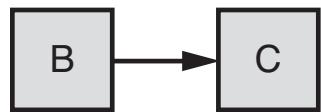
---

- $b$  is a belief state: a set of states
  - $o$  is an observation; a percept
  - $a$  is an action
- 
- $b' = \text{UPDATE}(\text{PREDICT}(b, a), o)$

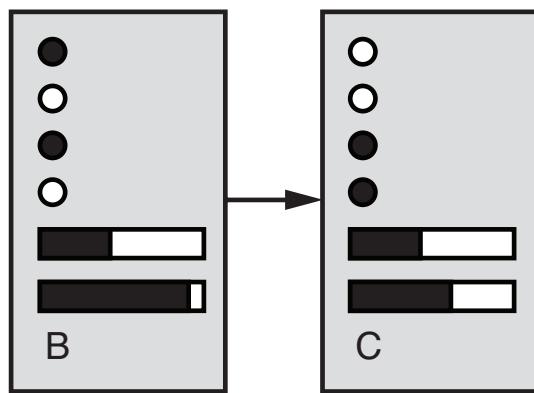
Kindergarten world: dirt may appear anywhere at any time,  
But actions are guaranteed to work.



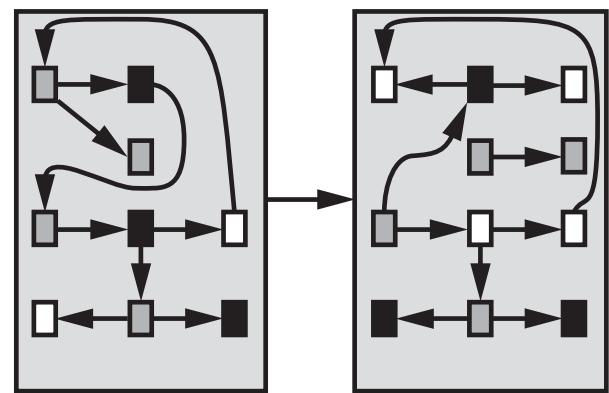
# State Representation



(a) Atomic

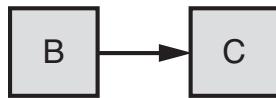


(b) Factored

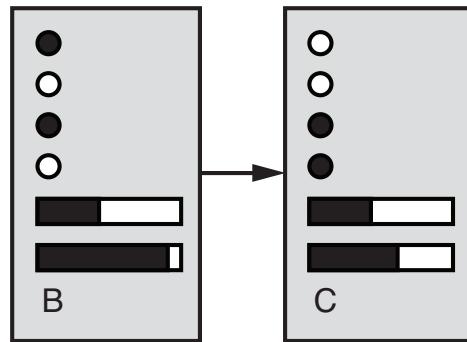


(b) Structured

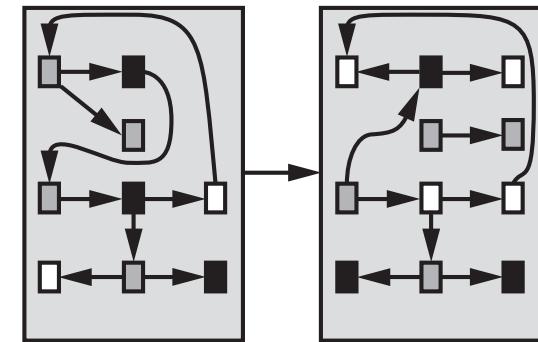
# State Representation



(a) Atomic



(b) Factored



(b) Structured

- ## Bayes Nets?

Factored

- ## SQL data base?

One table: Factored;

Several: Structured

- ## Java program?

Structured

# Representation

---

- Atomic Representation
  - $s_1 = s_2$ ;  $\text{Result}(s, a) = s'$ ;  $\text{GoalTest}(s)$ ;
- Factored Representation
  - $\text{Attribute}(s) = val \dots$  ( $val$  numeric or Boolean)
  - Result and GoalTest in terms of attributes
- Structured Representation
  - All of above
  - Relations, Functions:  $\text{Rel}(a, b, c)$ ;  $F(a, b) = c$
  - Objects; with Quantifiers  
(for all  $\forall x$ , there exists  $\exists y$ )

# Planning with Factored States

---

- World is made up of states which are defined in terms of state variables
  - Can be Boolean or categorical or continuous
  - Where have we used this before?
- State: *complete assignment over state variables*  
So,  $k$  Boolean state variables represent how many states?
- Actions change the values of the state variables
  - Applicability conditions of actions are also specified in terms of partial assignments over state variables

# “Classical” Planning

---

- **State:** conjunction of Boolean state variables
- **Action Schema:**
  - $\text{Action}(\text{Fly}(p, \text{from}, \text{to}),$   
Precond:  $\text{At}(p, \text{from}) \wedge \text{Plane}(p)$   
 $\wedge \text{Airport}(\text{from}) \wedge \text{Airport}(\text{to})$
  - Effect:  $\neg\text{At}(p, \text{from}) \wedge \text{At}(p, \text{to}))$

Implicitly defines  $\text{Actions}(s)$  and  $\text{Result}(s, a)$

# Expressiveness of the language

```
Init(At(C1, SFO) ∧ At(C2, JFK) ∧ At(P1, SFO) ∧ At(P2, JFK)
     ∧ Cargo(C1) ∧ Cargo(C2) ∧ Plane(P1) ∧ Plane(P2)
     ∧ Airport(JFK) ∧ Airport(SFO))
Goal(At(C1, JFK) ∧ At(C2, SFO))
Action(Load(c, p, a),
       PRECOND: At(c, a) ∧ At(p, a) ∧ Cargo(c) ∧ Plane(p) ∧ Airport(a)
       EFFECT: ¬At(c, a) ∧ In(c, p))
Action(Unload(c, p, a),
       PRECOND: In(c, p) ∧ At(p, a) ∧ Cargo(c) ∧ Plane(p) ∧ Airport(a)
       EFFECT: At(c, a) ∧ ¬In(c, p))
Action(Fly(p, from, to),
       PRECOND: At(p, from) ∧ Plane(p) ∧ Airport(from) ∧ Airport(to)
       EFFECT: ¬At(p, from) ∧ At(p, to))
```

**Figure 10.1** A PDDL description of an air cargo transportation planning problem.

# Advantages of the Language

---

- Natural to read, write, verify
- Abstract over similar actions
- Easy to extend with more complex syntax and semantics
- Can be compiled or interpreted
- Can automatically derived heuristics (relaxed problem)

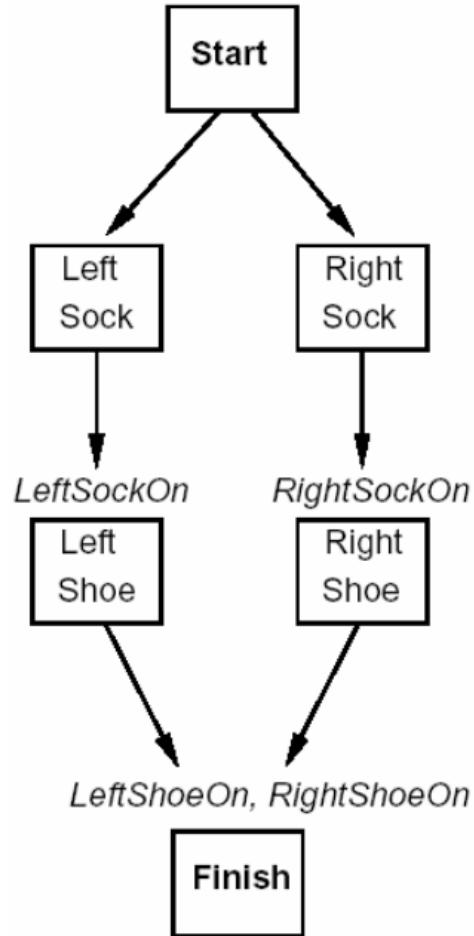
# Planning Algorithms

---

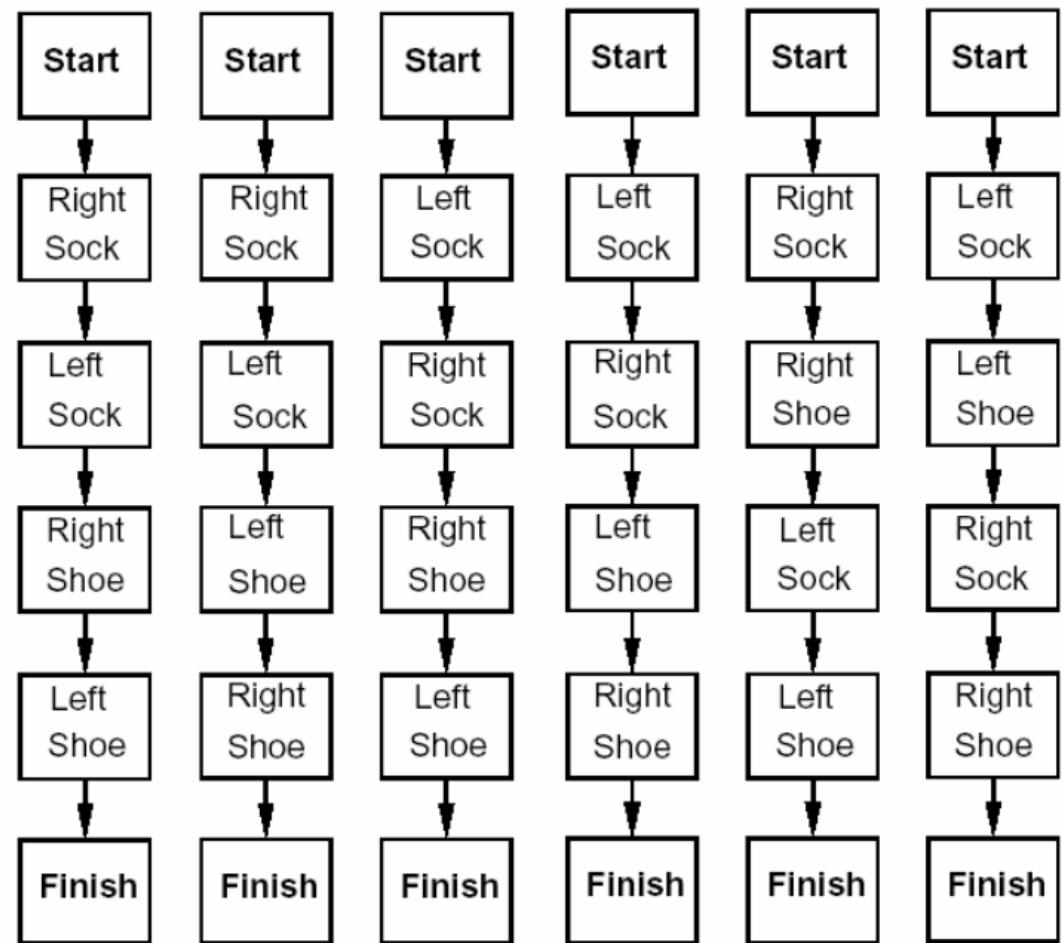
- Forward (progression) state-space search
  - ... it's just search
- Backward (regression) state-space search
  - Consider Goal:  $\text{Own}(0136042597)$   
 $\text{Action}(\text{Buy}(i), \text{Pre: } \text{ISBN}(i) \text{ Eff: } \text{Own}(i))$
  - In general, may involve unbound variables
- Plan-space search
  - Start with empty plan, add branches

# Plan-space search

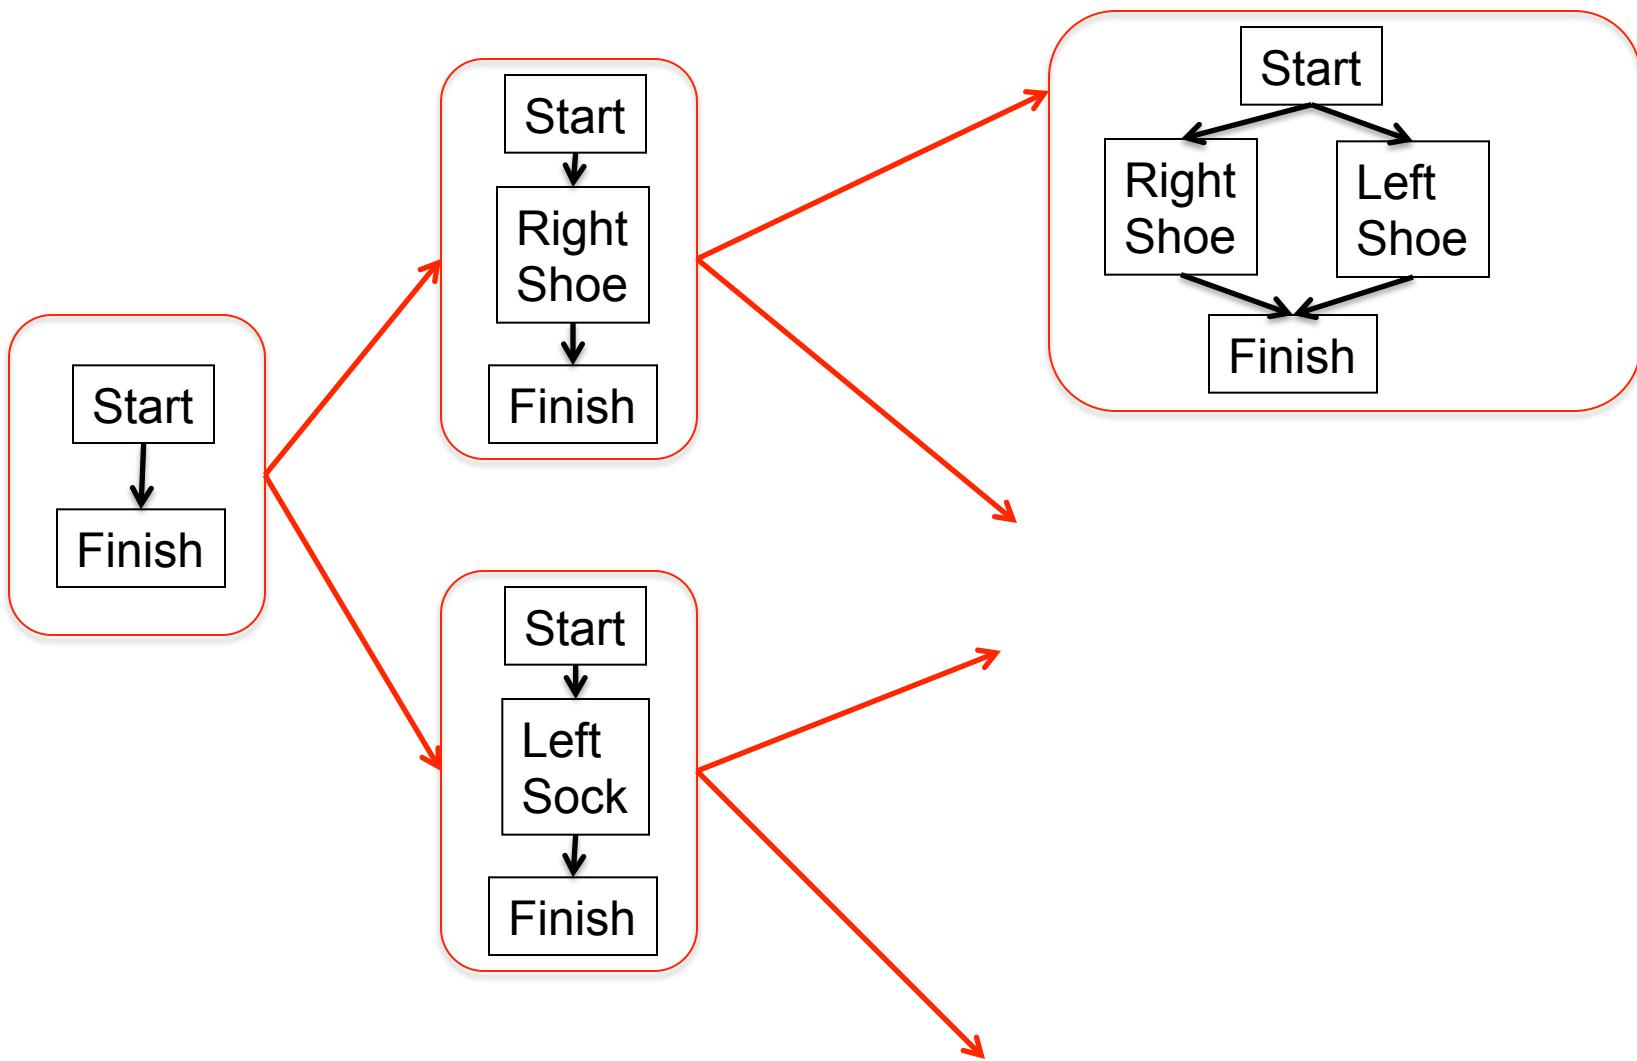
Partial Order Plan:



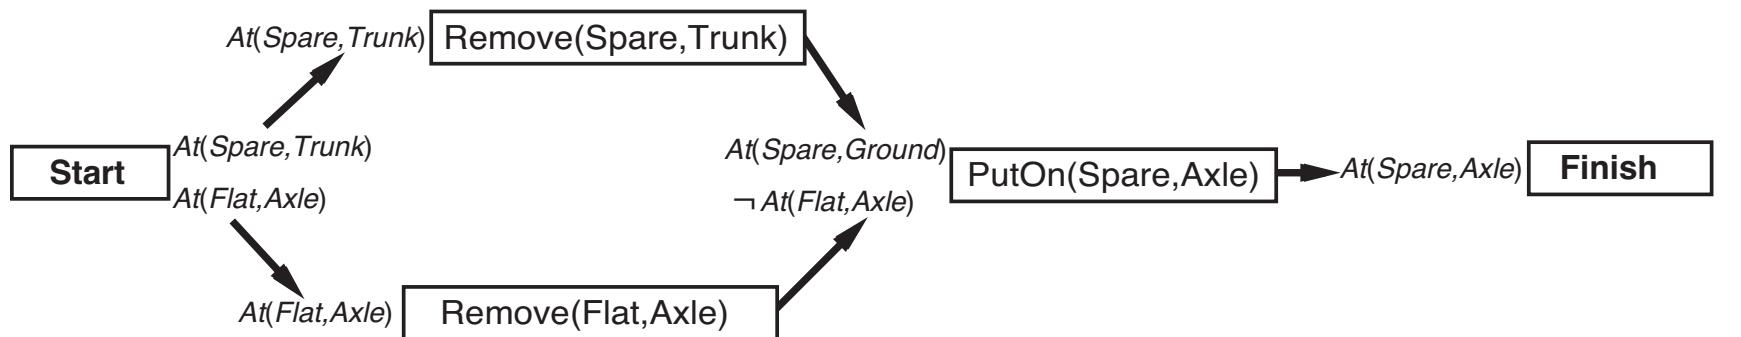
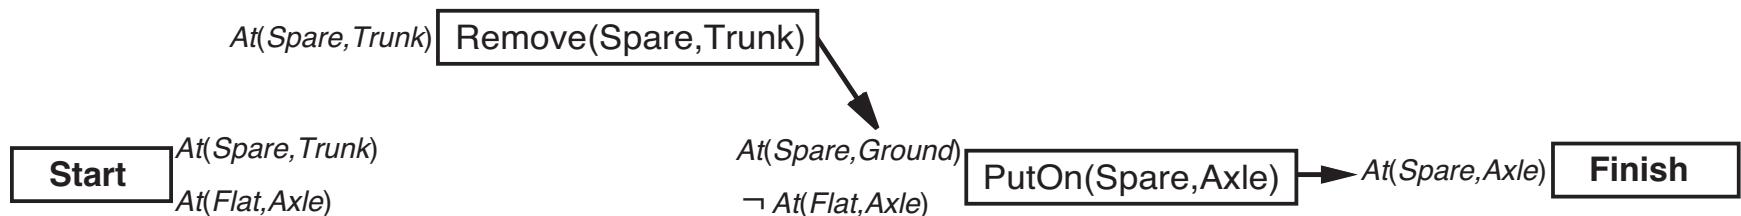
Total Order Plans:



# Plan-space search

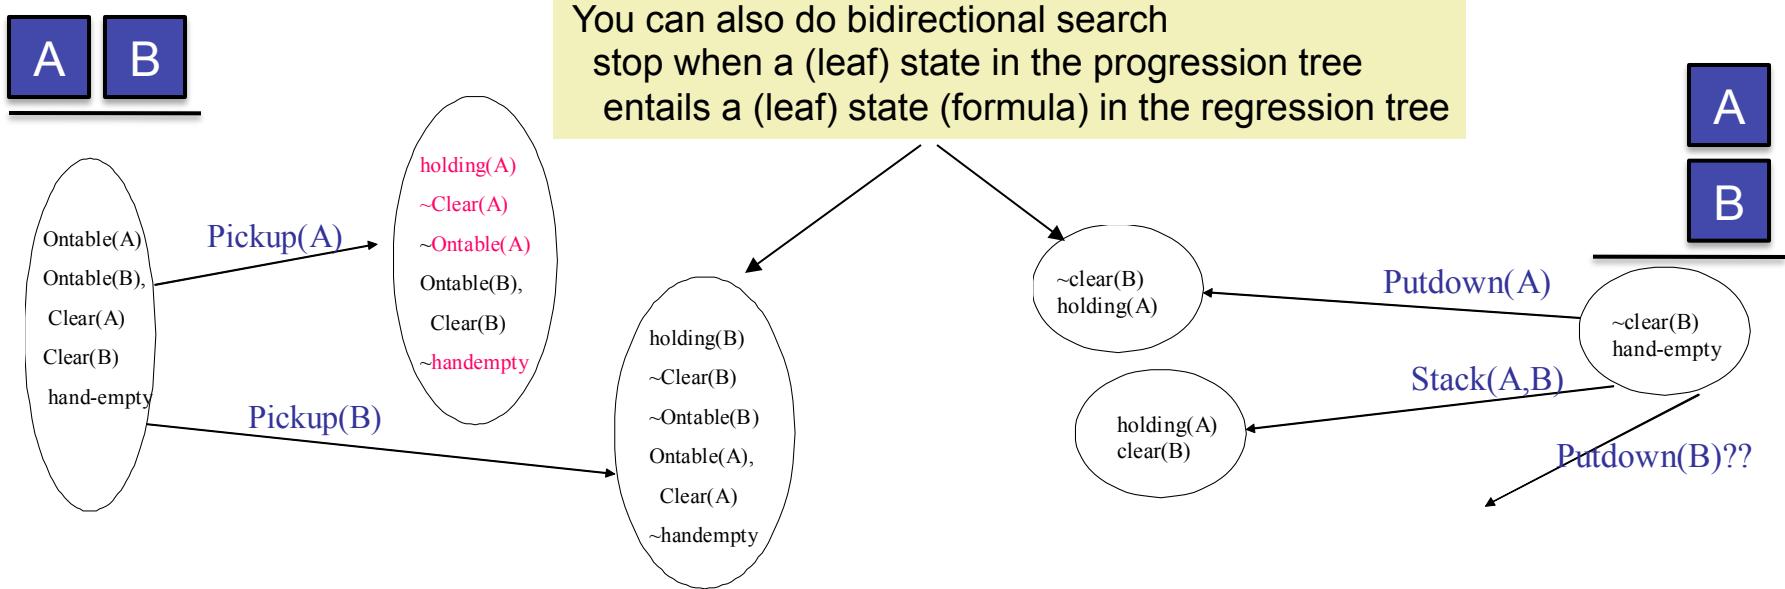


# Plan-space search



# Progression vs. Regression

- Progression has higher branching factor
- Progression searches in the space of complete (and consistent) states
- Regression has lower branching factor
- Regression searches in the space of *partial states*
  - There are  $3^n$  partial states (as against  $2^n$  complete)



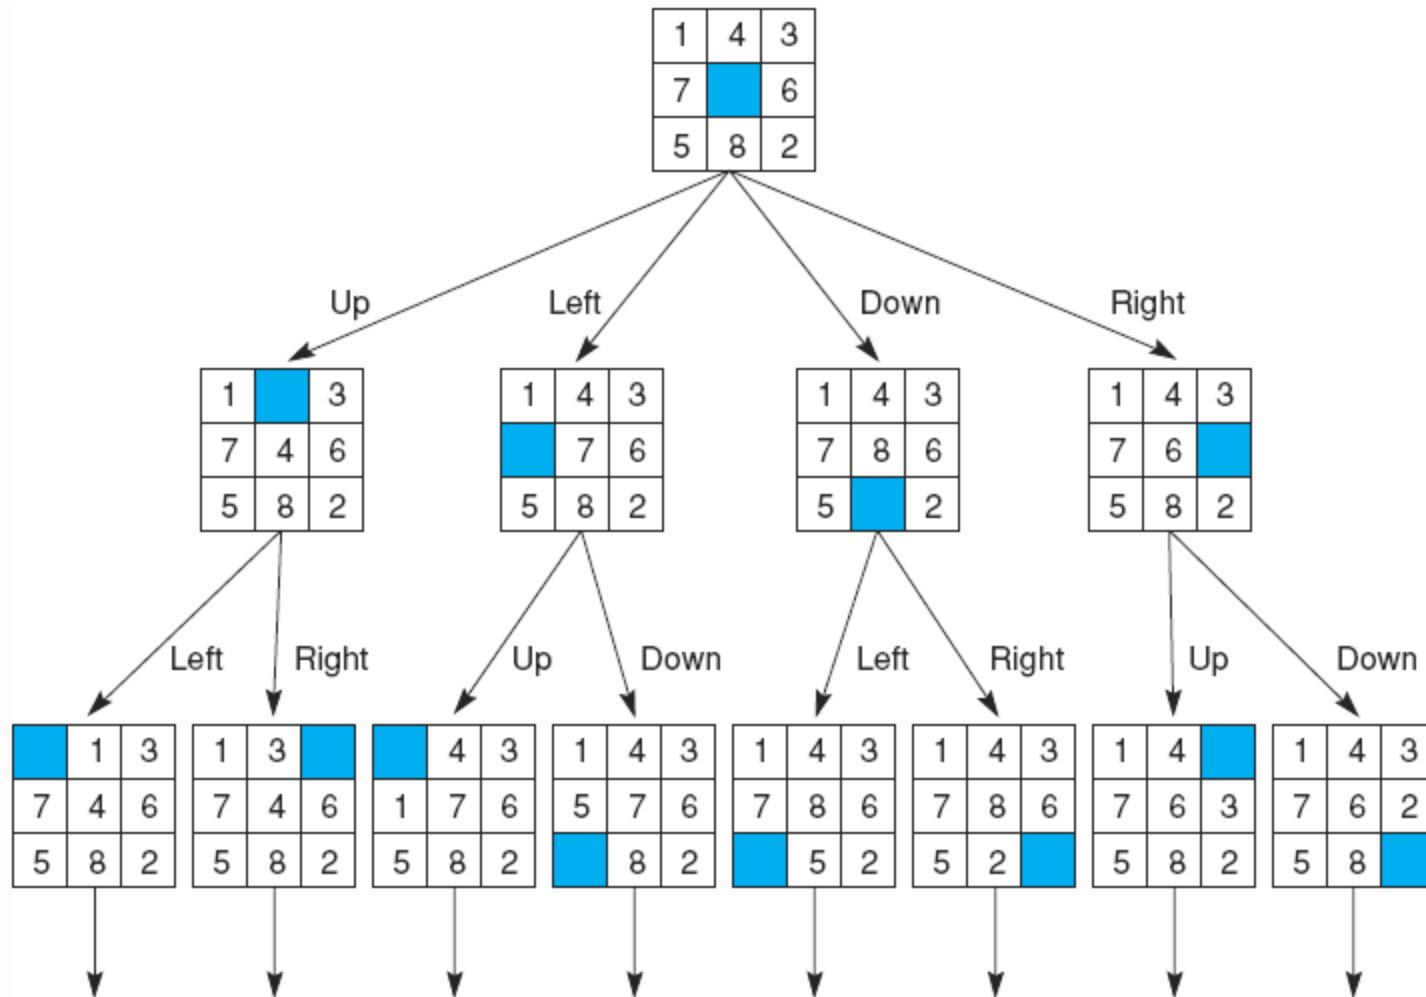
# State of the art

---

- Annual planning competitions
- Best technique has varied over time
- Currently: mostly forward state-space
- Largely due to good heuristics (relaxed prob.)
  - Heuristics for atomic (state search) problem  
Can only come from *outside* analysis of domain
  - Heuristics for factored (planning) problem  
Can be *domain-independent*

# 8-puzzle state space

---



# 8-puzzle action schema

---

- *Action(Slide(t, a, b),*

Pre:  $On(t, a) \wedge Tile(t) \wedge Blank(b) \wedge Adjacent(a, b)$

Eff:  $On(t, b) \wedge Blank(a) \wedge \neg On(t, a) \wedge \neg Blank(b))$

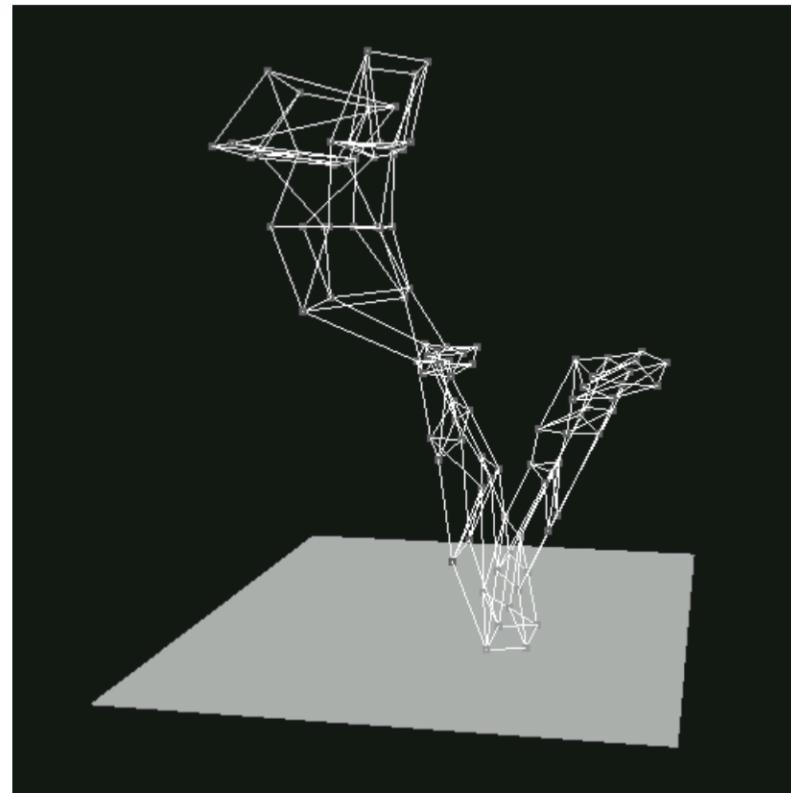
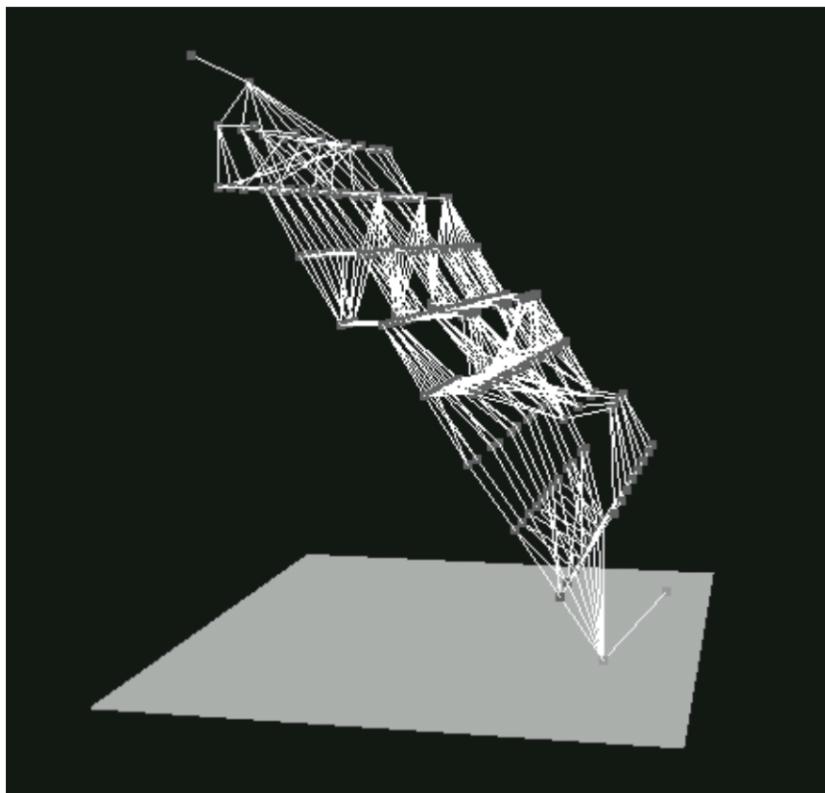
# 8-puzzle heuristics

---

- *Action(Slide( $t$ ,  $a$ ,  $b$ ),*  
Pre:  $On(t, a) \wedge Tile(t) \wedge Blank(b) \wedge Adjacent(a, b)$   
Eff:  $On(t, b) \wedge Blank(a) \wedge \neg On(t, a) \wedge \neg Blank(b)$ )
- Relaxed problems: ignore preconditions
  - $Blank(b)$ :
  - Manhattan distance heuristic
  - $Blank(b) \wedge Adjacent(a, b)$ :
  - Number-of-misplaced-tiles heuristic
- Relaxed problems: ignore delete lists
  - Ignore negative effects  
Convex space: Find solution (non-optimal) by hill-climbing

# Convex search: ignore del lists

---



# Factored Rep allows control

---

## Is there any Need for Domain-Dependent Control Information?

Matthew L. Ginsberg\* and Donald F. Geddis†

Computer Science Department

Stanford University

Stanford, California 94305

[ginsberg@cs.stanford.edu](mailto:ginsberg@cs.stanford.edu)

# Factored Rep allows control

---

## Is there any Need for Domain-Dependent Control Information?

Matthew L. Ginsberg\* and Donald F. Geddis†

Computer Science Department

Stanford University

Stanford, California 94305

[ginsberg@cs.stanford.edu](mailto:ginsberg@cs.stanford.edu)

### Abstract

No.

### 1 Introduction

The split between *base-level* and *metalevel* knowledge has long been recognized by the declarative community. Roughly speaking, base-level knowledge has to do with information about some particular domain, while metalevel knowledge has to do with knowledge *about* that information. A typical base-level fact might be, "Iraq invaded Kuwait," while a typical metalevel fact might be, "To show that a country *c* is aggressive, first try to find another country that has been invaded by *c*."

because sentences expressing modal knowledge typically involve the use of predicate symbols that have other sentences as arguments. Thus a typical modal sentence might be, "I know that Iraq invaded Kuwait," or "I don't know of anyone that Kuwait has invaded." Note that the information here doesn't refer to the domain so much as it does to our knowledge *about* the domain; nor is it control information telling us what to do with this knowledge. It simply reports on the state of our information at some particular point in time. In general, we will describe as modal all information describing the structure of our declarative knowledge in some way.

The claim we are making – and we intend to prove it – is that there is no place in declarative systems

# Beyond Classical Planning

---

- Convenient to have more expressive lang.  
“Move *all* the cargo from SFO to JFK”
- Can be done with careful extensions to factored planning language
- Or: Use existing first-order logical provers
- Strong foundation for studying planning
- Still, less used in practice than other techniques

# First-Order Logic

---

- And, Or, Not, Implies  
(as in propositional logic)
- Variables ranging over *objects*
- Relations and functions over objects
- Quantifiers  $\forall$  (for all) and  $\exists$  (exists)
  - Goal:  $\forall c \text{ Cargo}(c) \Rightarrow \text{At}(c, \text{JFK})$

# Situation Calculus

---

- Actions are objects
- Situations are objects
- Function:  $s_2 = \text{Result}(s, a)$
- Fluents:  $\text{At}(C_1, \text{JFK}, s)$  *change over time*
- Possibility axioms
  - Say when an action is possible
- Successor-state axioms
  - Partially describe the resulting state of an action

# Situation Calculus

---

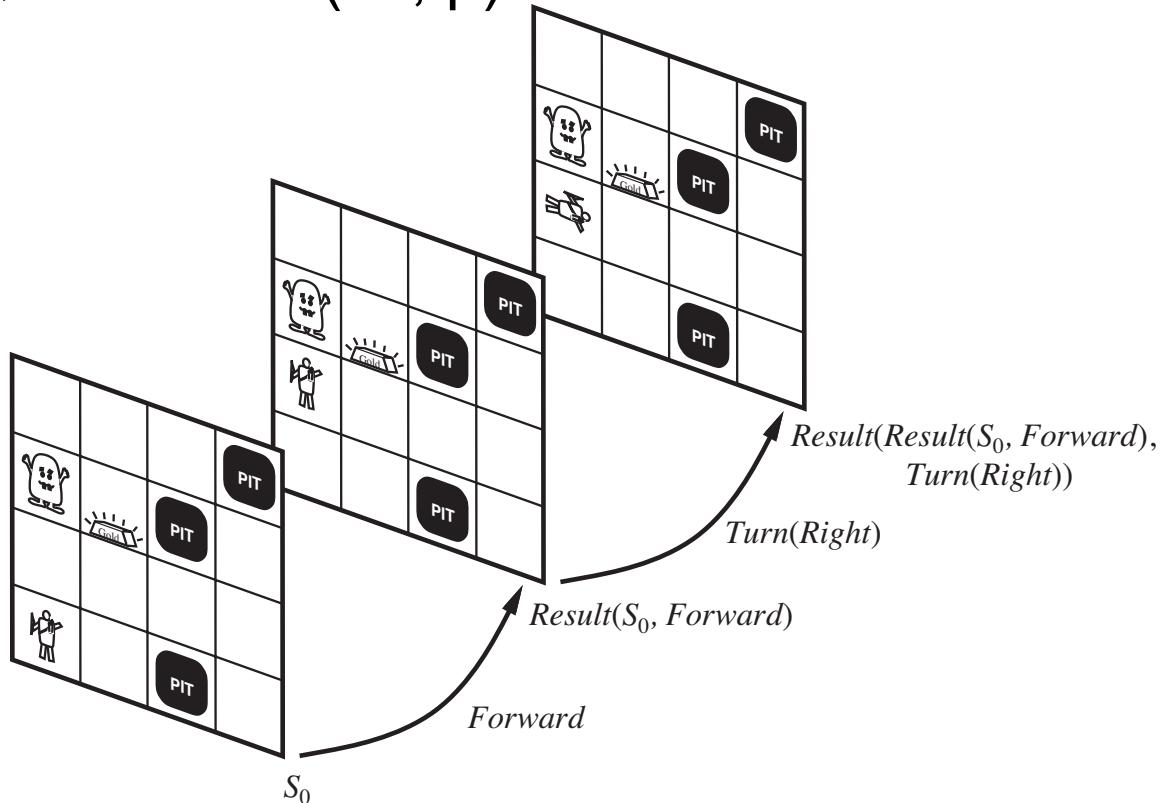
- Possibility Axioms (for each action)
  - $\text{SomeFormula}(s) \Rightarrow \text{Poss}(a, s)$
  - $\text{Alive}(\text{Agent}, s) \wedge \text{Have}(\text{Agent}, \text{Arrow}, s) \Rightarrow \text{Poss}(\text{Shoot}, s)$
- Successor-state Axiom (for each fluent)
  - $\text{Poss}(a, s) \Rightarrow (\text{fluent is true} \Leftrightarrow a \text{ made it true} \vee \text{it was true and } a \text{ left it alone})$
  - $\text{Poss}(a, s) \Rightarrow (\text{Holding}(\text{Agent}, g, \text{Result}(s, a)) \Leftrightarrow a = \text{Grab}(g) \vee (\text{Holding}(\text{Agent}, g, s) \wedge a \neq \text{Release}(g)))$

# Situations as Result of Action

$\exists s, p : \text{Goal}(s) \wedge s = \text{Result}(s_0, p)$

Situation Calculus

First-order Logic



$s = \text{Result}(s, [])$

$\text{Result}(s, [a, b, \dots]) = \text{Result}(\text{Result}(s, a), [b, \dots])$

# Planning Graphs

---

- Planning graphs are an efficient way to create a representation of a planning problem that can be used to
  - Achieve better heuristic estimates
  - Directly construct plans
- Planning graphs only work for propositional problems
  - Compile to propositional if necessary

# Planning Graphs

---

- Planning graphs consists of a seq of levels that correspond to time steps in the plan.
  - Level 0 is the initial state.
  - Each level consists of a set of literals and a set of actions that represent what *might* be possible at that step in the plan
  - *Might be* is the key to efficiency
  - Records only a restricted subset of possible negative interactions among actions.

# Planning Graphs

---

- Each level consists of
- **Literals** = all those that *could* be true at that time step, depending upon the actions executed at preceding time steps.
- **Actions** = all those actions that *could* have their preconditions satisfied at that time step, depending on which of the literals actually hold.

# Planning Graph Example

---

Init(Have(Cake))

Goal(Have(Cake)  $\wedge$  Eaten(Cake))

Action(Eat(Cake),

    PRECOND: Have(Cake)

    EFFECT:  $\neg$ Have(Cake)  $\wedge$  Eaten(Cake))

Action(Bake(Cake),

    PRECOND:  $\neg$  Have(Cake)

    EFFECT: Have(Cake))

# Planning Graph Example

---

$S_0$

$A_0$

$S_1$

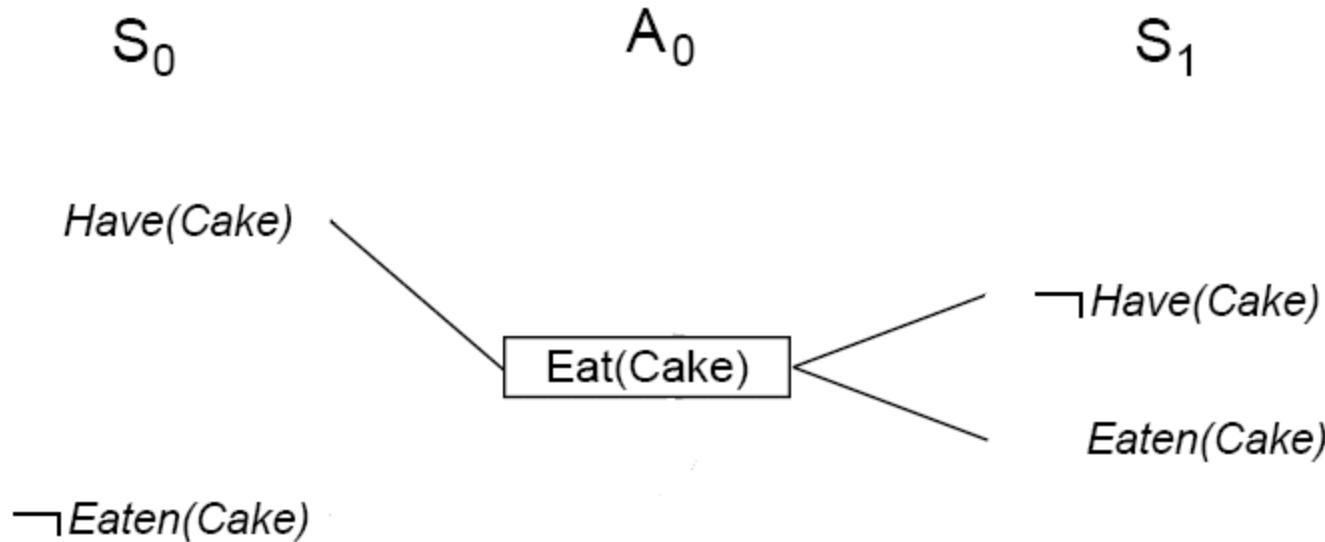
*Have(Cake)*

$\neg Eaten(Cake)$

Create level 0 from initial problem state.

# Planning Graph Example

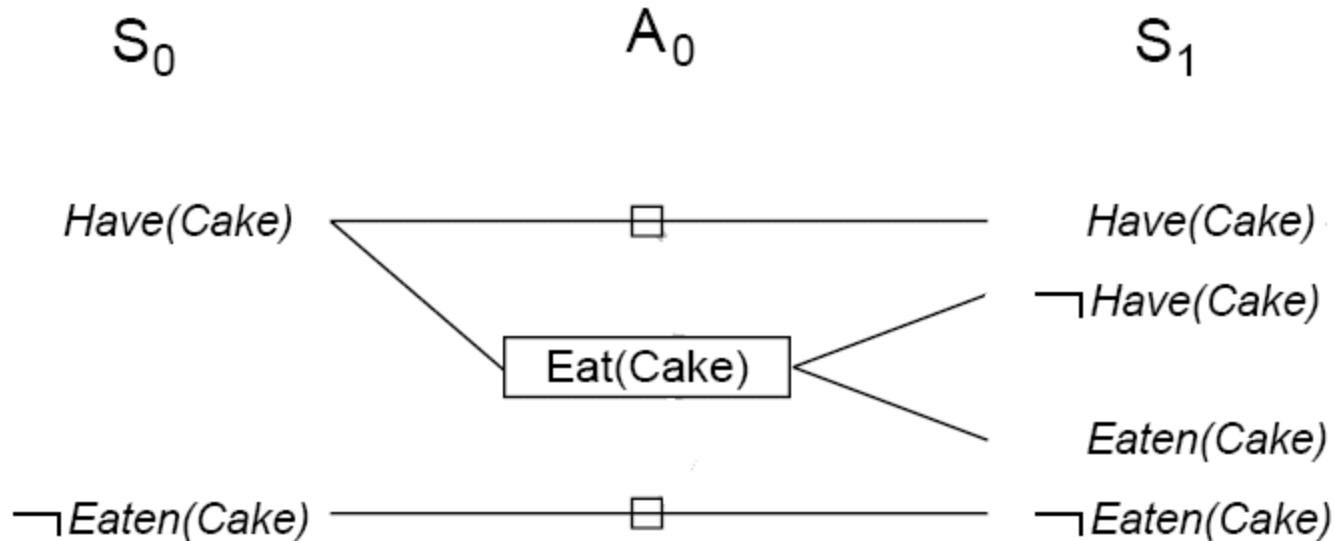
---



Add all applicable actions.

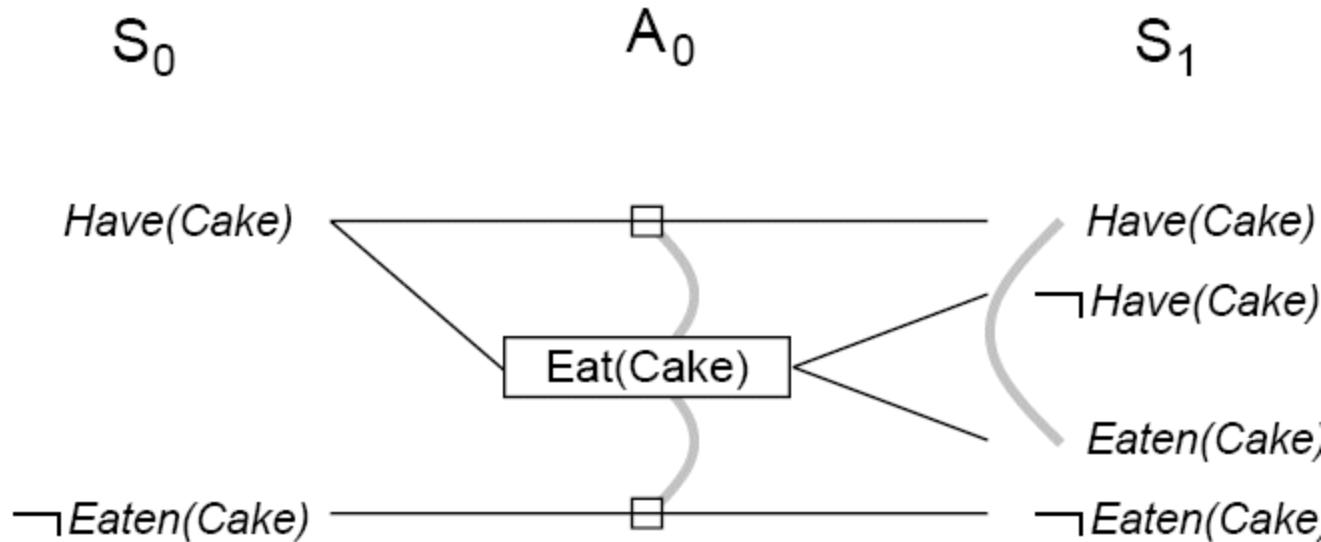
Add all effects to the next state.

# Planning Graph Example



Add *persistence actions* (inaction = no-ops) to map all literals in state  $S_i$  to state  $S_{i+1}$ .

# Planning Graph Example



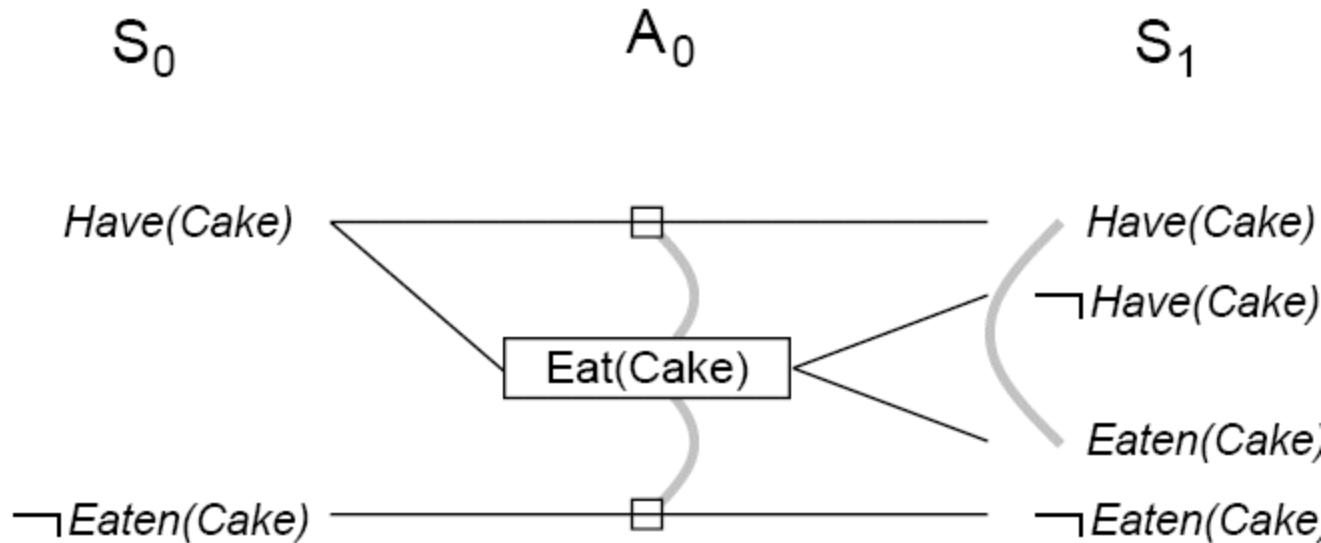
Identify *mutual exclusions* between actions and literals based on potential conflicts.

# Mutual exclusion

---

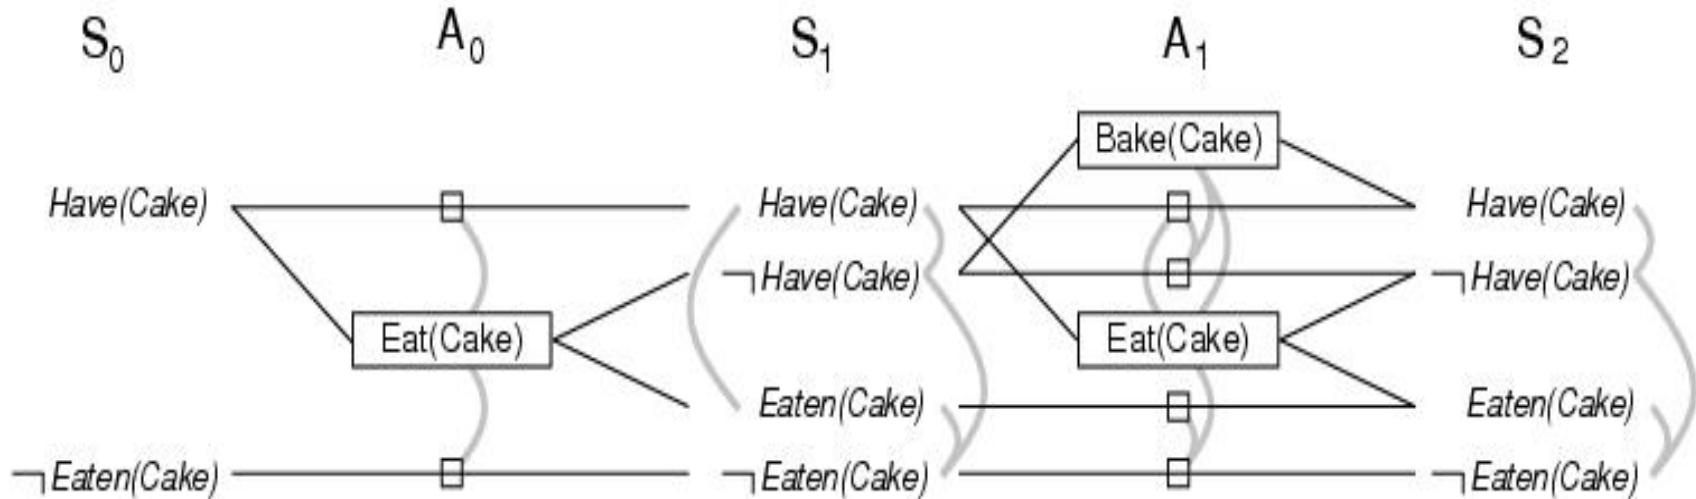
- A mutex relation holds between **two actions** when:
  - *Inconsistent effects*: one action negates the effect of another.
  - *Interference*: one of the effects of one action is the negation of a precondition of the other.
  - *Competing needs*: one of the preconditions of one action is mutually exclusive with the precondition of the other.
- A mutex relation holds between **two literals** when:
  - one is the negation of the other OR
  - each possible action pair that could achieve the literals is mutex (inconsistent support).

# Cake example



- Level  $S_1$  contains all literals that could result from picking any subset of actions in  $A_0$ 
  - Conflicts between literals that can not occur together (as a consequence of the selection action) are represented by mutex links.
  - $S_1$  defines multiple states and the mutex links are the constraints that define this set of states.

# Cake example



- Repeat process until graph levels off:
  - two consecutive levels are identical

# PG and Heuristic Estimation

---

- PG's provide information about the problem
  - PG is a relaxed problem.
  - A literal that does not appear in the final level of the graph cannot be achieved by any plan.
    - $h(s) = \infty$
  - Level Cost: First level in which a goal appears
    - Very low estimate, since several actions can occur
    - Improvement: restrict to one action per level using *serial PG* (add mutex links between every pair of actions, except persistence actions).

# PG and Heuristic Estimation

---

- Cost of a conjunction of goals
  - Max-level: maximum first level of any of the goals
  - Sum-level: sum of first levels of all the goals
  - Set-level: First level in which all goals appear without being mutex

# The GRAPHPLAN Algorithm

---

- Extract a solution directly from the PG

**function** GRAPHPLAN(*problem*) **return** *solution* or failure

*graph*  $\leftarrow$  INITIAL-PLANNING-GRAFH(*problem*)

*goals*  $\leftarrow$  GOALS[*problem*]

**loop do**

**if** *goals* all non-mutex in last level of *graph* **then do**

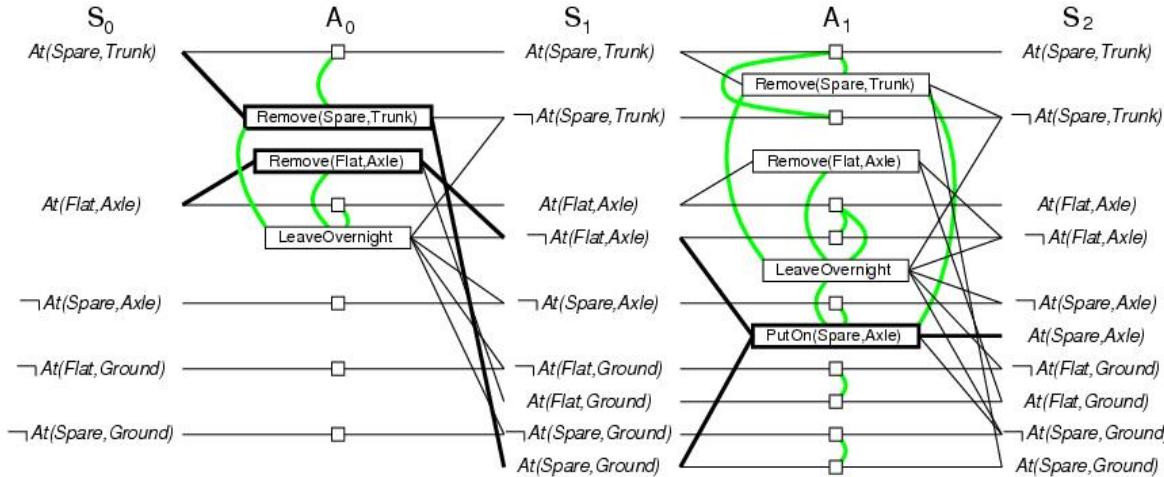
*solution*  $\leftarrow$  EXTRACT-SOLUTION(*graph*, *goals*, LEN(*graph*))

**if** *solution*  $\neq$  failure **then return** *solution*

**else if** NO-SOLUTION-POSSIBLE(*graph*) **then return** failure

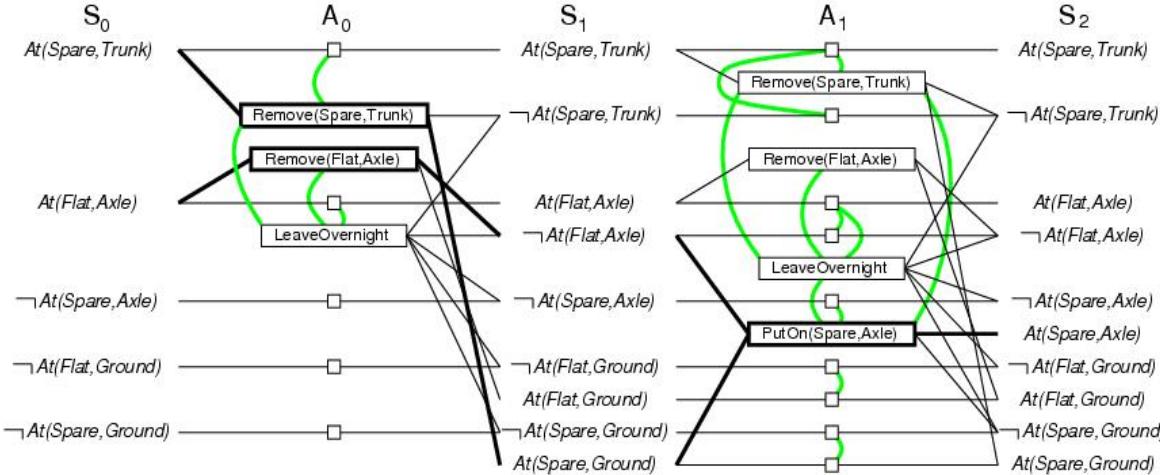
*graph*  $\leftarrow$  EXPAND-GRAFH(*graph*, *problem*)

# GRAPHPLAN example



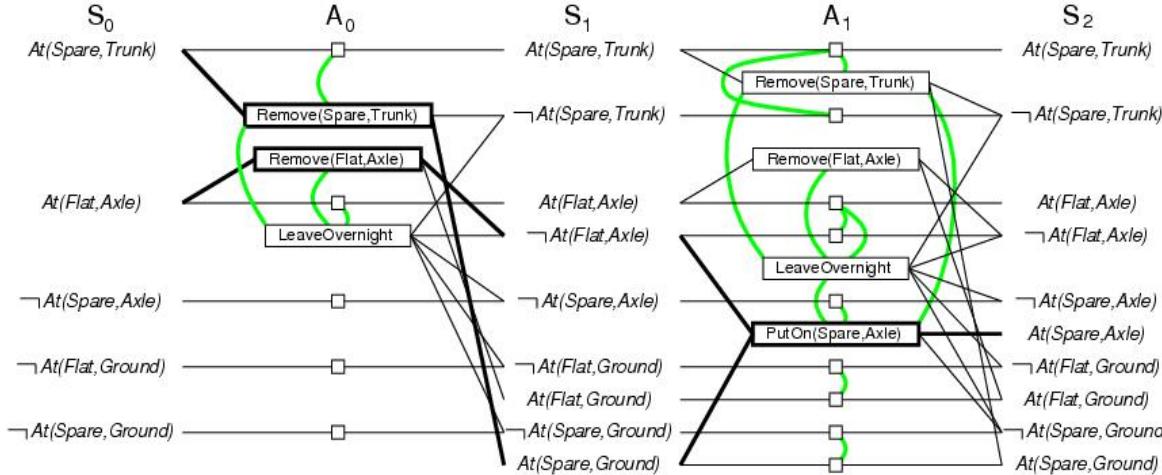
- Initially the plan consist of 5 literals from the initial state (S<sub>0</sub>).
- Add actions whose preconditions are satisfied by EXPAND-GRAFH (A<sub>0</sub>)
- Also add persistence actions and mutex relations.
- Add the effects at level S<sub>1</sub>
- Repeat until goal is in level S<sub>i</sub>

# GRAPHPLAN example



- EXPAND-GRAF also looks for mutex relations
  - Inconsistent effects
    - E.g.  $Remove(Spare, Trunk)$  and  $LeaveOverNight$  due to  $At(Spare, Ground)$  and **not**  $At(Spare, Ground)$
  - Interference
    - E.g.  $Remove(Flat, Axle)$  and  $LeaveOverNight$   $At(Flat, Axle)$  as PRECOND and **not**  $At(Flat, Axle)$  as EFFECT
  - Competing needs
    - E.g.  $PutOn(Spare, Axle)$  and  $Remove(Flat, Axle)$  due to  $At(Flat, Axle)$  and **not**  $At(Flat, Axle)$
  - Inconsistent support
    - E.g. in  $S_2$ ,  $At(Spare, Axle)$  and  $At(Flat, Axle)$

# GRAPHPLAN example



- In  $S_2$ , the goal literals exist and are not mutex with any other
  - Solution might exist and EXTRACT-SOLUTION will try to find it
- EXTRACT-SOLUTION can search with:
  - Initial state = last level of PG and goal goals of planning problem
  - Actions = select any set of non-conflicting actions that cover the goals in the state
  - Goal = reach level  $S_0$  such that all goals are satisfied
  - Cost = 1 for each action.

# GRAPHPLAN Termination

---

- Termination of graph construction? YES
- PG are monotonically increasing or decreasing:
  - Literals increase monotonically
  - Actions increase monotonically
  - Mutexes decrease monotonically
- Because of these properties and because there is a finite number of actions and literals, every PG will eventually level off