

# CS 221: Artificial Intelligence

## Lecture 8: MDPs

Sebastian Thrun and Peter Norvig

Slide credit: Dan Klein, Stuart Russell

# Rhino Museum Tourguide

---



# Minerva Robot

---



# Pearl Robot

---

## Nursebot Pearl

Cocktail Hour at the  
Longwood Independent  
Living Facility

# Mine Mapping Robot Groundhog

---





# Mine Mapping Robot Groundhog

---












# Planning: Classical Situation


---



- World deterministic
- State observable

# MDP-Style Planning

---




- Policy
- Universal Plan
- Navigation function

- World stochastic
- State observable

# Stochastic, Partially Observable



---



[Sondik 72] [Littman/Cassandra/Kaelbling 97]


# Stochastic, Partially Observable

---




# Stochastic, Partially Observable

---



# Stochastic, Partially Observable

---



# A Quiz

---

# states	sensors	actions	size belief space?
3	perfect	deterministic	3: $s_1, s_2, s_3$
3	perfect	stochastic	3: $s_1, s_2, s_3$
3	abstract states	deterministic	$2^3-1$ : $s_1, s_2, s_3, s_{12}, s_{13}, s_{23}, s_{123}$
3	stochastic	deterministic	2-dim continuous: $p(S=s_1), p(S=s_2)$
3	none	stochastic	2-dim continuous: $p(S=s_1), p(S=s_2)$
1-dim continuous	stochastic	deterministic	$\infty$ -dim continuous
1-dim continuous	stochastic	stochastic	$\infty$ -dim continuous
$\infty$ -dim continuous	stochastic	stochastic	aargh!





# MPD Planning

---

- Solution for Planning problem
  - Noisy controls
  - Perfect perception
  - Generates “universal plan” (=policy)


# Grid World

- The agent lives in a grid
- Walls block the agent's path
- The agent's actions do not always go as planned:
  - 80% of the time, the action North takes the agent North (if there is no wall there)
  - 10% of the time, North takes the agent West; 10% East
  - If there is a wall in the direction the agent would have been taken, the agent stays put
- Small “living” reward each step
- Big rewards come at the end
- Goal: maximize sum of rewards\*




# Grid Futures

## Deterministic Grid World





## Stochastic Grid World



# Markov Decision Processes


- An MDP is defined by:
  - A **set of states**  $s \in S$
  - A **set of actions**  $a \in A$
  - A **transition function**  $T(s, a, s')$ 
    - Prob that  $a$  from  $s$  leads to  $s'$
    - i.e.,  $P(s' | s, a)$
    - Also called the model
  - A **reward function**  $R(s, a, s')$ 
    - Sometimes just  $R(s)$  or  $R(s')$
  - A **start state** (or distribution)
  - Maybe a **terminal state**
- MDPs are a family of non-deterministic search problems
  - Reinforcement learning: MDPs where we don't know the transition or reward functions



# Solving MDPs


- In deterministic single-agent search problems, want an optimal **plan**, or sequence of actions, from start to a goal
- In an MDP, we want an optimal **policy**  $\pi^*: S \rightarrow A$ 
  - A policy  $\pi$  gives an action for each state
  - An optimal policy maximizes expected utility if followed
  - Defines a reflex agent

Optimal policy when  
 $R(s, a, s') = -0.03$  for  
all non-terminals  $s$




# Example Optimal Policies


---




$$R(s) = -0.01$$



$$R(s) = -0.03$$




$$R(s) = -0.4$$



$$R(s) = -2.0$$

# MDP Search Trees


- Each MDP state gives a search tree






# Why Not Search Trees?

- Why not solve with conventional planning?
- Problems:
  - This tree is usually infinite (why?)
  - Same states appear over and over (why?)
  - We would search once per state (why?)



# Utilities




- Utility = sum of future reward
- Problem: infinite state sequences have infinite rewards
- Solutions:
  - Finite horizon:
    - Terminate episodes after a fixed T steps (e.g. life)
    - Gives nonstationary policies ( $\pi$  depends on time left)
  - Absorbing state: guarantee that for every policy, a terminal state will eventually be reached (like “done” for High-Low)
  - Discounting: for  $0 < \gamma < 1$

$$U([r_0, \dots, r_\infty]) = \sum_{t=0}^{\infty} \gamma^t r_t \leq R_{\max} / (1 - \gamma)$$

- Smaller  $\gamma$  means smaller “horizon” – shorter term focus

# Discounting

- Typically discount rewards by  $\gamma < 1$  each time step
  - Sooner rewards have higher utility than later rewards
  - Also helps the algorithms converge




# Recap: Defining MDPs

---

- Markov decision processes:

- States  $S$
- Start state  $s_0$
- Actions  $A$
- Transitions  $P(s' | s, a)$  (or  $T(s, a, s')$ )
- Rewards  $R(s, a, s')$  (and discount  $\gamma$ )




- MDP quantities so far:

- Policy = Choice of action for each state
- Utility (or return) = sum of discounted rewards

# Optimal Utilities

- Fundamental operation: compute the values (optimal expect-max utilities) of states  $s$
- Why? Optimal values define optimal policies!
- Define the value of a state  $s$ :  
 $V^*(s)$  = expected utility starting in  $s$  and acting optimally
- Define the value of a q-state  $(s,a)$ :  
 $Q^*(s,a)$  = expected utility starting in  $s$ , taking action  $a$  and thereafter acting optimally
- Define the optimal policy:  
 $\pi^*(s)$  = optimal action from state  $s$




3	0.812	0.868	0.912	<span style="border: 1px solid black; padding: 2px;">+1</span>
2	0.762		0.660	<span style="border: 1px solid black; padding: 2px;">-1</span>
1	0.705	0.655	0.611	0.388
	1	2	3	4

3	→	→	→	<span style="border: 1px solid black; padding: 2px;">+1</span>
2	↑		↑	<span style="border: 1px solid black; padding: 2px;">-1</span>
1	↑	←	←	←
	1	2	3	4

# The Bellman Equations

- Definition of “optimal utility” leads to a simple one-step lookahead relationship amongst optimal utility values:

Optimal rewards = maximize over first action and then follow optimal policy



- Formally:

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$


# Solving MDPs

- We want to find the **optimal policy**  $\pi^*$
- Proposal 1: modified expect-max search, starting from each state  $s$ :

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a Q^*(s, a)$$





# Value Iteration

---


- Idea:

- Start with  $V_0^*(s) = 0$
- Given  $V_i^*$ , calculate the values for all states for depth  $i+1$ :

$$V_{i+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$

- This is called a **value update** or **Bellman update**
  - Repeat until convergence
- Theorem: will converge to unique optimal values
    - Basic idea: approximations get refined towards optimal values
    - Policy may converge long before values do

# Example: Bellman Updates



$$V_{i+1}(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$

$$V_2(\langle 3, 3 \rangle) = \sum_{s'} T(\langle 3, 3 \rangle, \text{right}, s') [R(\langle 3, 3 \rangle) + 0.9 V_1(s')]$$

max happens for  
 $a=\text{right}$ , other  
actions not shown

$$= 0.9 [0.8 \cdot 1 + 0.1 \cdot 0 + 0.1 \cdot 0]$$

# Example: Value Iteration

---

$V_2$

3	0	0	0.72	<span style="border: 1px solid black; padding: 2px;">+1</span>
2	0		0	<span style="border: 1px solid black; padding: 2px;">-1</span>
1	0	0	0	0
	1	2	3	4

$V_3$

3	0	0.52	0.78	<span style="border: 1px solid black; padding: 2px;">+1</span>
2	0		0.43	<span style="border: 1px solid black; padding: 2px;">-1</span>
1	0	0	0	0
	1	2	3	4

- Information propagates outward from terminal states and eventually all states have correct value estimates

# Computing Actions

---

- Which action should we chose from state  $s$ :
  - Given optimal values  $V$ ?

$$\arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Given optimal q-values  $Q$ ?

$$\arg \max_a Q^*(s, a)$$

- Lesson: actions are easier to select from  $Q$ 's!



# Asynchronous Value Iteration\*

---


- In value iteration, we update every state in each iteration
- Actually, *any* sequences of Bellman updates will converge if every state is visited infinitely often
- In fact, we can update the policy as seldom or often as we like, and we will still converge
- Idea: Update states whose value we expect to change:  
If  $|V_{i+1}(s) - V_i(s)|$  is large then update predecessors of  $s$

# Value Iteration: Example


---



# Another Example




Map




Value Function and Plan



# Another Example




Map



Value Function and Plan

# Stochastic, Partially Observable

---






# Value Iteration in Belief space: POMDPs




---

- Partially Observable Markov Decision Process
  - Known model (learning even harder!)
  - Observation uncertainty
  - Usually also: transition uncertainty
  - Planning in belief space = space of all probability distributions
- Value function: Piecewise linear, convex function over the belief space



# Introduction to POMDPs (1 of 3)




# Introduction to POMDPs (2 of 3)



# Introduction to POMDPs (3 of 3)



$$V(p(s_1)) = \sum_{z=\{A,B\}} V(p(s_1 | z)) p(z)$$



# POMDP Algorithm

---

- Belief space = Space of all probability distribution (continuous)
- Value function: Max of set of linear functions in belief space
- Backup: Create new linear functions
- Number of linear functions can grow fast!

# Why is This So Complex?

---

State Space Planning  
(no state uncertainty)



Belief Space Planning  
(full state uncertainties)





# Belief Space Structure

---


The controller may be globally uncertain...

but not usually.



# Augmented MDPs:


---



# Path Planning with Augmented MDPs

---

Conventional planner



Probabilistic Planner

