

CS 221: Artificial Intelligence

Reinforcement Learning

Peter Norvig and Sebastian Thrun

Slide credit: Dan Klein, Stuart Russell, Andrew Moore



Review: Markov Decision Processes

- An MDP is defined by:

- set of states $s \in S$
- set of actions $a \in A$
- transition function $P(s' | s, a)$
aka $T(s, a, s')$
- reward function $R(s, a, s')$
or $R(s, a)$ or $R(s')$

- Problem:

- Find policy $\pi(s)$ to maximize $\sum \gamma^t R(s, \pi(s), s')$



Value Iteration

- Idea:

- Start with $U_0(s) = 0$
- Given $U_i(s)$, calculate:

$$U_{i+1}(s) \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U_i(s')$$

- Repeat until convergence

- Theorem: will converge to unique optimal utilities

$$U^\pi(s) = R(s) + \gamma \sum_{s'} P(s' | s, \pi(s)) U^\pi(s') .$$

- But: Requires model (R and P)

3				<div>+1</div>
2				<div>-1</div>
1	START			
	1	2	3	4


Reinforcement Learning

- What if the state transition function P and the reward function R are *unknown*?



Example: Backgammon

- Reward only for win / loss in terminal states, zero otherwise
- TD-Gammon (Tesauro, 1992) learns a function approximation to $U(s)$ using a neural network
- Combined with depth 3 search, one of the top 3 human players in the world – good at positional judgment, but mistakes in end game beyond depth 3



Example: Animal Learning

- RL studied experimentally for more than 60 years in psychology
 - Rewards: food, pain, hunger, drugs, etc.
 - Mechanisms and sophistication debated
- Example: foraging
 - Bees learn near-optimal foraging plan in field of artificial flowers with controlled nectar supplies
 - Bees have a direct neural connection from nectar intake measurement to motor planning area



What to Learn

- If you don't know P , R , what you learn depends on what you want to do
 - **Utility-based agent:** learn $U(s)$
 - **Q-learning agent:** learn utility $Q(s, a)$
 - **Reflex agent:** learn policy $\pi(s)$
- And on how patient/reckless you are
 - **Passive:** use a fixed policy π
 - **Active:** modify policy as you go

Passive Temporal-Difference

function PASSIVE-TD-AGENT(*percept*) **returns** an action

inputs: *percept*, a percept indicating the current state s' and reward signal r'

persistent: π , a fixed policy

U , a table of utilities, initially empty

N_s , a table of frequencies for states, initially zero

s, a, r , the previous state, action, and reward, initially null

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

 increment $N_s[s]$

$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$

if $s'.\text{TERMINAL?}$ **then** $s, a, r \leftarrow \text{null}$ **else** $s, a, r \leftarrow s', \pi[s'], r'$


return a

Example

0	0	0	<div>+1</div>
0		0	<div>-1</div>
0	0	0	0


Example

$$U^\pi(s) \leftarrow U^\pi(s) + \alpha(R(s) + \gamma U^\pi(s') - U^\pi(s))$$




Example

$$U^\pi(s) \leftarrow U^\pi(s) + \alpha(R(s) + \gamma U^\pi(s') - U^\pi(s))$$



Example

$$U^\pi(s) \leftarrow U^\pi(s) + \alpha(R(s) + \gamma U^\pi(s') - U^\pi(s))$$



Example

$$U^\pi(s) \leftarrow U^\pi(s) + \alpha(R(s) + \gamma U^\pi(s') - U^\pi(s))$$

0	0.8	0.92	<div>+1</div>
0		0	<div>-1</div>
0	0	0	0

Example

$$U^\pi(s) \leftarrow U^\pi(s) + \alpha(R(s) + \gamma U^\pi(s') - U^\pi(s))$$

.20	.28	.36	<div>+1</div>
.17		-.12	<div>-1</div>
.12	-0.01	-.16	-.2

Sample results




Figure 21.5 The TD learning curves for the 4×3 world. (a) The utility estimates for a selected subset of states, as a function of the number of trials. (b) The root-mean-square error in the estimate for $U(1, 1)$, averaged over 20 runs of 500 trials each. Only the first 100 trials are shown to enable comparison with Figure 21.3.

Problem

- The problem
 - Algorithm computes the value of one specific policy
 - May never try the best moves
 - May never visit some states
 - May waste time computing unneeded utilities


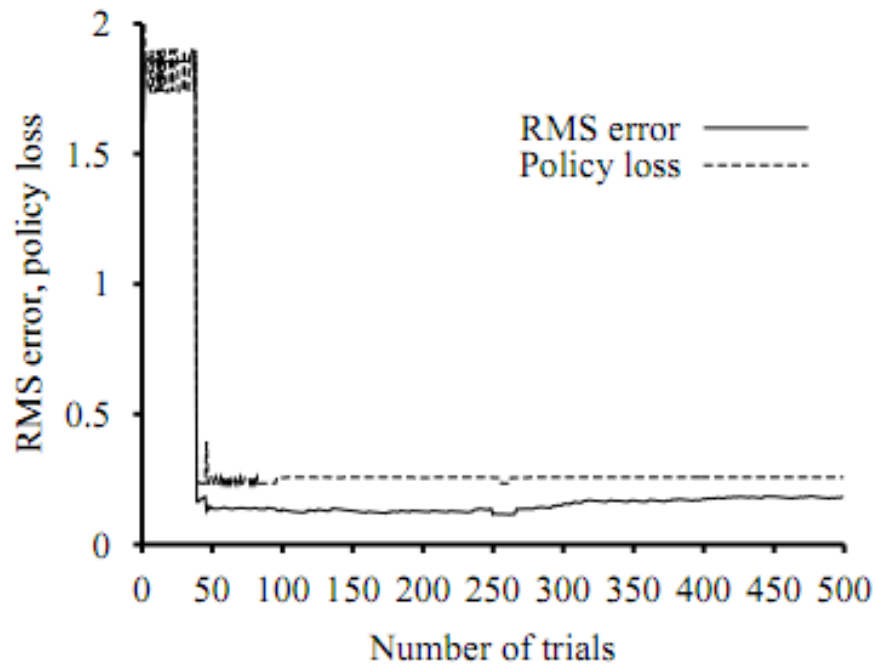
Quiz

- Is there a fix?
- Can we explore different policies (active reinforcement learning), yet still learn the optimal policy?

Greedy TD-Agent

- Learn U with TD-Learning
- After each change to U , solve MDP to compute new optimal policy $\pi(s)$
- Continue learning with new policy



Greedy Agent Results



■ Wrong policy

- Too much exploitation, not enough exploration
- Fix: higher $U(s)$ for unexplored states

Exploratory Agent



- Converges to correct policy

Alternative: Q-Learning


- Instead of learning $U(s)$, learn $Q(a,s)$: the utility of action a in state s
- With utilities, need a model, P , to act:

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U(s')$$


- With Q-values, no model needed
 - Compute Q-values with TD/exploration algorithm


$$Q(s, a) \leftarrow Q(s, a) + \alpha (R(s) + \gamma Q(s', a') - Q(s, a))$$

Q-Learning




Q-Learning





Q-Learning




$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma Q(s', a') - Q(s, a))$$

Q-Learning

<div>0000</div>	<div>00.330</div>	<div>00.780</div>	<div>+1</div>
<div>0000</div>	<div></div>	<div>0000</div>	<div>-1</div>
<div>0000</div>	<div>0000</div>	<div>0000</div>	<div>0000</div>



Q-Learning

- Q-learning produces tables of q-values:



Q-Learning Properties

- Amazing result: Q-learning converges to optimal policy
 - If you explore enough
 - If you make the learning rate small enough
 - ... but not decrease it too quickly!
 - Basically doesn't matter how you select actions (!)
- Neat property: off-policy learning
 - learn optimal policy without following it (some caveats, see SARSA)






Practical Q-Learning

- In realistic situations, we cannot possibly learn about every single state!
 - Too many states to visit them all in training
 - Too many states to hold the q-tables in memory
- Instead, we want to generalize:
 - Learn about some small number of training states from experience
 - Generalize that experience to new, similar states
 - This is a fundamental idea in machine learning, and we'll see it over and over again


Generalization Example

- Let's say we discover through experience that this state is bad:
- In naïve Q learning, we know nothing about this state or its Q states:
- Or even this one!



Feature-Based Representations

- Solution: describe a state using a vector of features
 - Features are functions from states to real numbers (often 0/1) that capture important properties of the state
 - Example features:
 - Distance to closest ghost
 - Distance to closest dot
 - Number of ghosts
 - $1 / (\text{dist to dot})^2$
 - Is Pacman in a tunnel? (0/1)
 - etc.
 - Can also describe a Q-state (s, a) with features (e.g. action moves closer to food)



Linear Feature Functions

- Using a feature representation, we can write a Q function (or Utility function) for any state using a few weights:

$$V(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- Advantage: our experience is summed up in a few powerful numbers
- Disadvantage: states may share features but be very different in value!

Function Approximation

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- Q-learning with linear Q-functions:

$$Q(s, a) \leftarrow Q(s, a) + \alpha [\text{error}]$$

$$w_i \leftarrow w_i + \alpha [\text{error}] f_i(s, a)$$

- Intuitive interpretation:
 - Adjust weights of active features
 - E.g. if something unexpectedly bad happens, disprefer all states with that state's features (prudence or superstition?)

Summary: MDPs and RL

Things we know how to do:

- If we know the MDP

- Compute U^* , Q^* , π^* exactly
- Evaluate a fixed policy π

- If we don't know the MDP

- We can estimate the MDP then solve
- We can estimate U for a fixed policy π
- We can estimate $Q^*(s,a)$ for the optimal policy while executing an exploration policy

Techniques:

- Model-based DP

- Value Iteration
- Policy iteration

- Model-based RL

- Model-free RL:

- Value learning
- Q-learning